AI assistance for innovation adoption

Data Capacity Excellence

Pankaj Acharya
DCE-Product Owner

July 16th, 2026

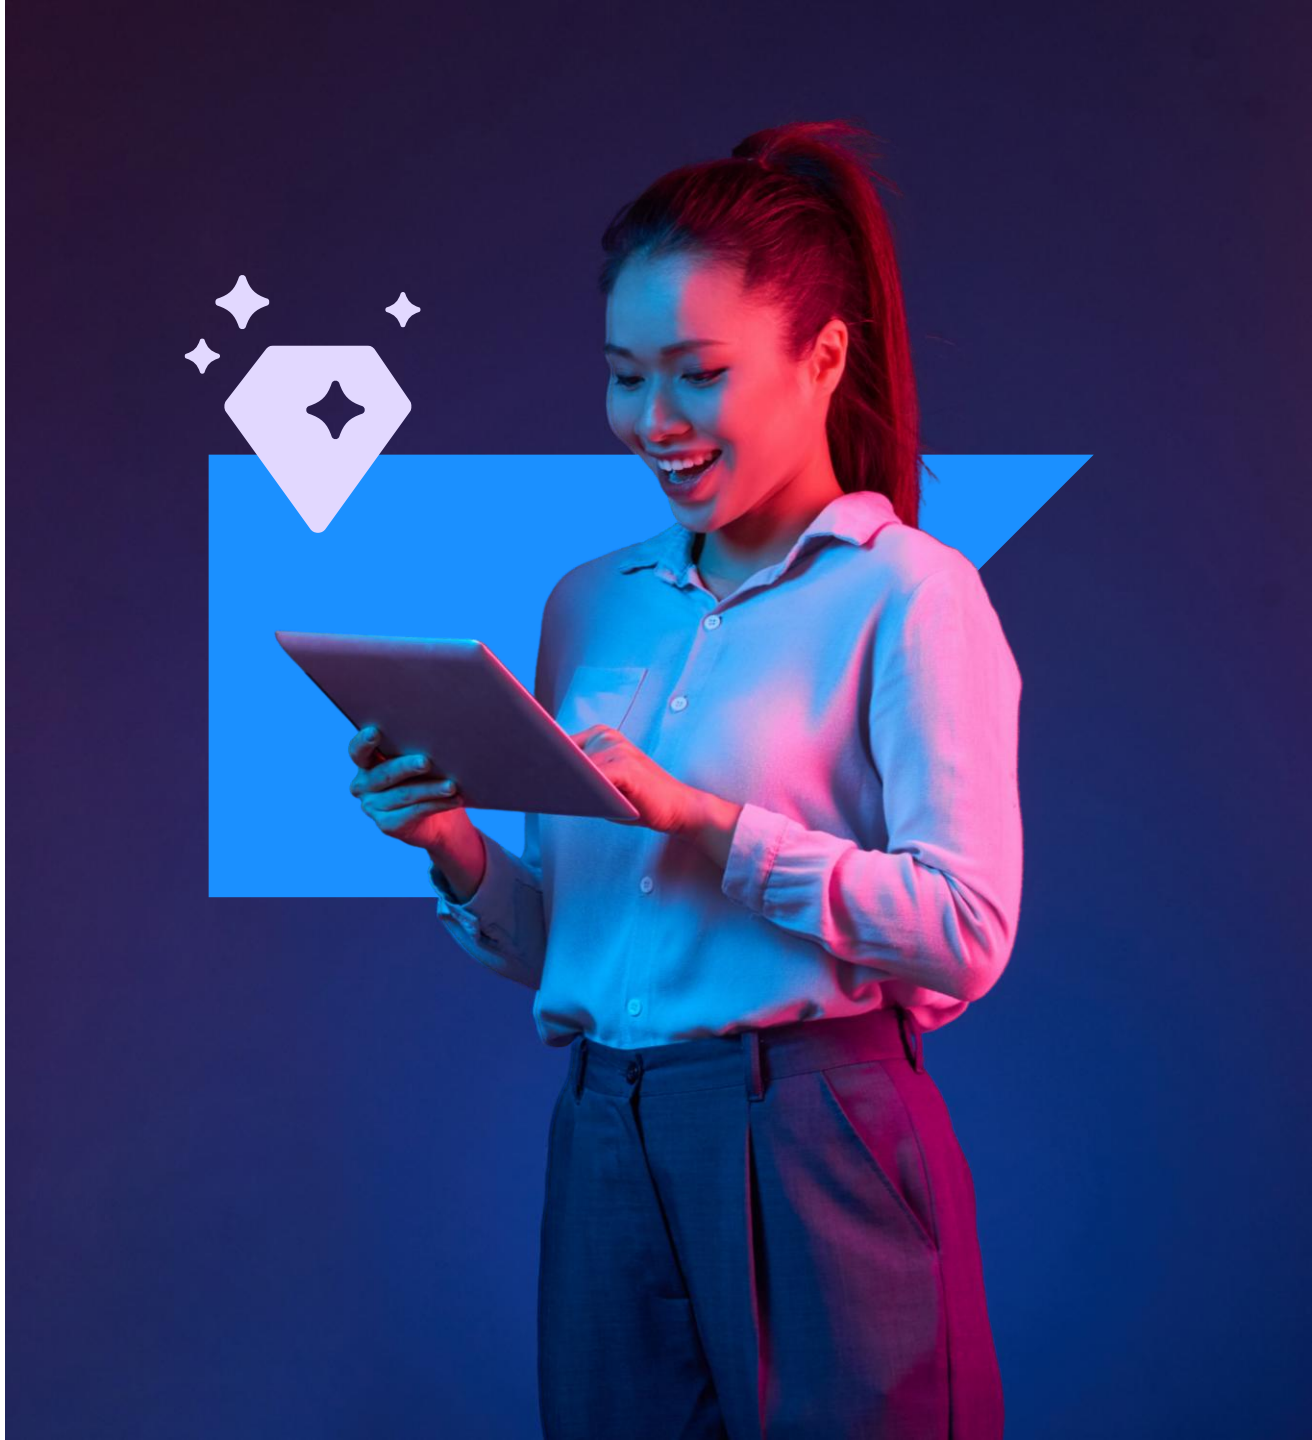

PUBLIC

Global Data Volume Management (DVM)

Our focus areas

- Global DVM service expert access to support you with only needed services
- Global DVM as RISE Transformation Partner
- Suite4Adoption as packaged & direct measurable E2E-Implementation Solution
- AI assistance for innovation adoption Data Capacity Excellence

About us

We are the Global Data Volume Management experts – a globally supporting and dedicated team of service developers and delivery experts helping customers take control of their data sizes and requirements across their IT landscapes.

Our team reduces data footprints – to decrease memory sizes – through archiving, deletion, summarization, index optimization, SAP Native Storage Extension (NSE), and co-creates DVM strategies that keep data accessible for audits, finance and analytics depending on your customer individual need.

Through our deep expertise and clear governance, the DVM service portfolio is shaped and grown to deliver clarity, predictability and resilience in data-volume growth to reduce unnecessary hardware costs and protecting system performance. We support you to reduce your IT TCO tremendously.

Customer Services & Delivery | Customer Engagement & Advisory | Global Delivery Hub | Cross-Solution Delivery | Application Integration

> DVM Suite4Adoption cuts IT Infrastructure costs by more than 25% .
>
> Steve Böttcher, BMW AG

Want to explore more?

Explore our **DVM Services** (ID: 9500320) @ SAP Services and Support Catalog

See our **DVM Cloud Services** (ID: 9507688) @ SAP Services and Support Catalog

Connect and engage with our **SAP Community**

Want to reach out?

Request a **DVM Scoping Expert**
m: DVM_Services_EMEA@sap.com

Tjark Riechers
Service Delivery Owner, Global DVM
m: tjark.riechers@sap.com

Lukas Diemand
Lead in Suite4Adoption, Global DVM
m: lukas.diemand@sap.com

André Gellert
Head of Global DVM, Global DVM
m: a.gellert@sap.com

Adoption is your key to value realization

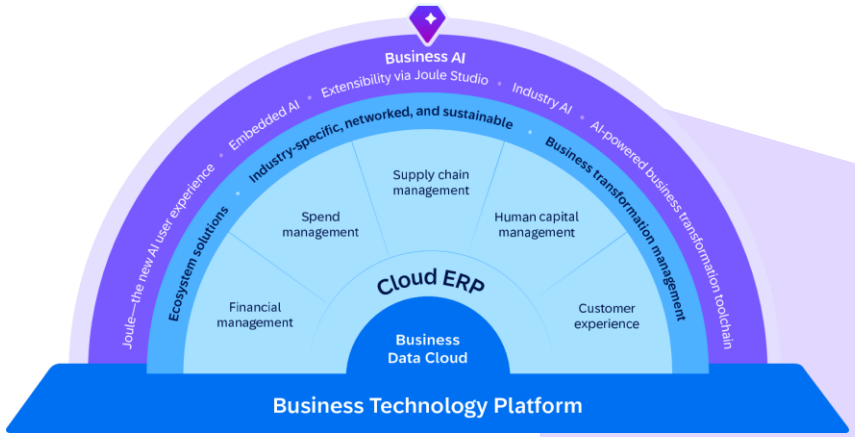

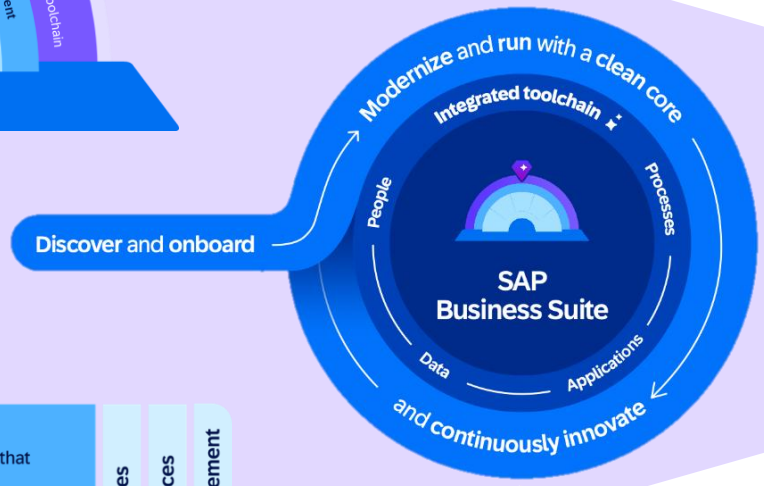

Applications & Technologies

\+ Integrated Toolchain

\+ Advanced and Max Success Plans

=

AI assistance for innovation adoption

It is a product, not a service

Productized expertise, designed to empower

To establish a process: “Continuous Improvement”

Runs in Cloud ALM, has a focus on cross-solution challenges

Success plans

Max | **Inclusive of all plans**
Enable large and complex business transformation with a strategic partnership that provides end-to-end guidance from dedicated experts.

Advanced | **Builds on the foundational plan**
Continually unlock the full value of your investment with proactive guidance as you deploy solutions and adopt innovations.

Foundational | **Included with every SAP cloud solution**
Adopt, transform, and ensure continuity with content, guidance, learning, mission-critical support, and more.

\+ Professional services
\+ Development services
\+ Application management

Agenda

1 AI assistance for Innovation Adoption in a Nutshell

2 Pre-Requisites

3 Content Overview

4 Data Capacity Excellence

5 Q & A

An intelligent, automated optimization engine or product to strengthen your SAP landscape

Collect Data

Collect raw data from the production system via
SAP Cloud ALM.

Generate Insights

Transform raw data into data-driven insights through AI-powered analysis and preprocessing.

Visualize Data

Gain a deep understanding of the data with a Service Report and a SAC Dashboard to visualize the data.

Articulate Issues

Convert identified insights into issues within SAP Cloud ALM
Issue and Action Management.

Create Guided Actions

Generate the recommended actions by SAP's Unique
Premium Knowledge.

Implement and track in PTM

Manage and monitor implementation through project
and setup (PTM) to ensure observability and quality.

Capability library to drive optimization & improvement

Data Capacity Excellence

Optimize data usage
and storage

Operational Excellence of SAP BTP

Improve platform
operations

Evolution Readiness

Prepare for future
SAP innovations

Enduring Cybersecurity

Align security with
business priorities

Technical Efficiency Optimization

Improve system performance
and stability

Adoption Potential for RISE with SAP

Take full advantage of
all solution entitlements

Extensibility Intelligence

Reduce technical debt and
maintain a cleaner core

AI assistance for innovation adoption

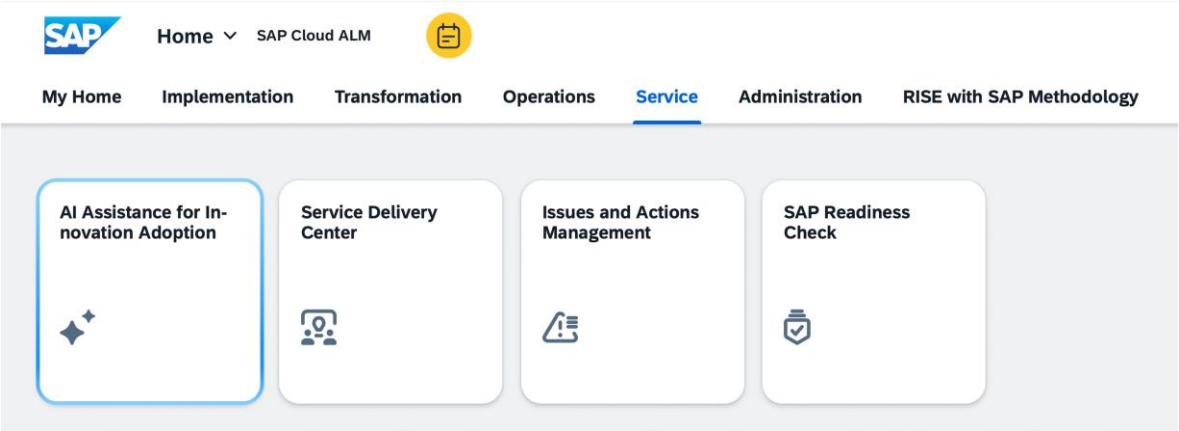

Built on SAP Cloud ALM, it combines SAP Analytics Cloud dashboards with AI-supported insights and guided actions.

Landing Page

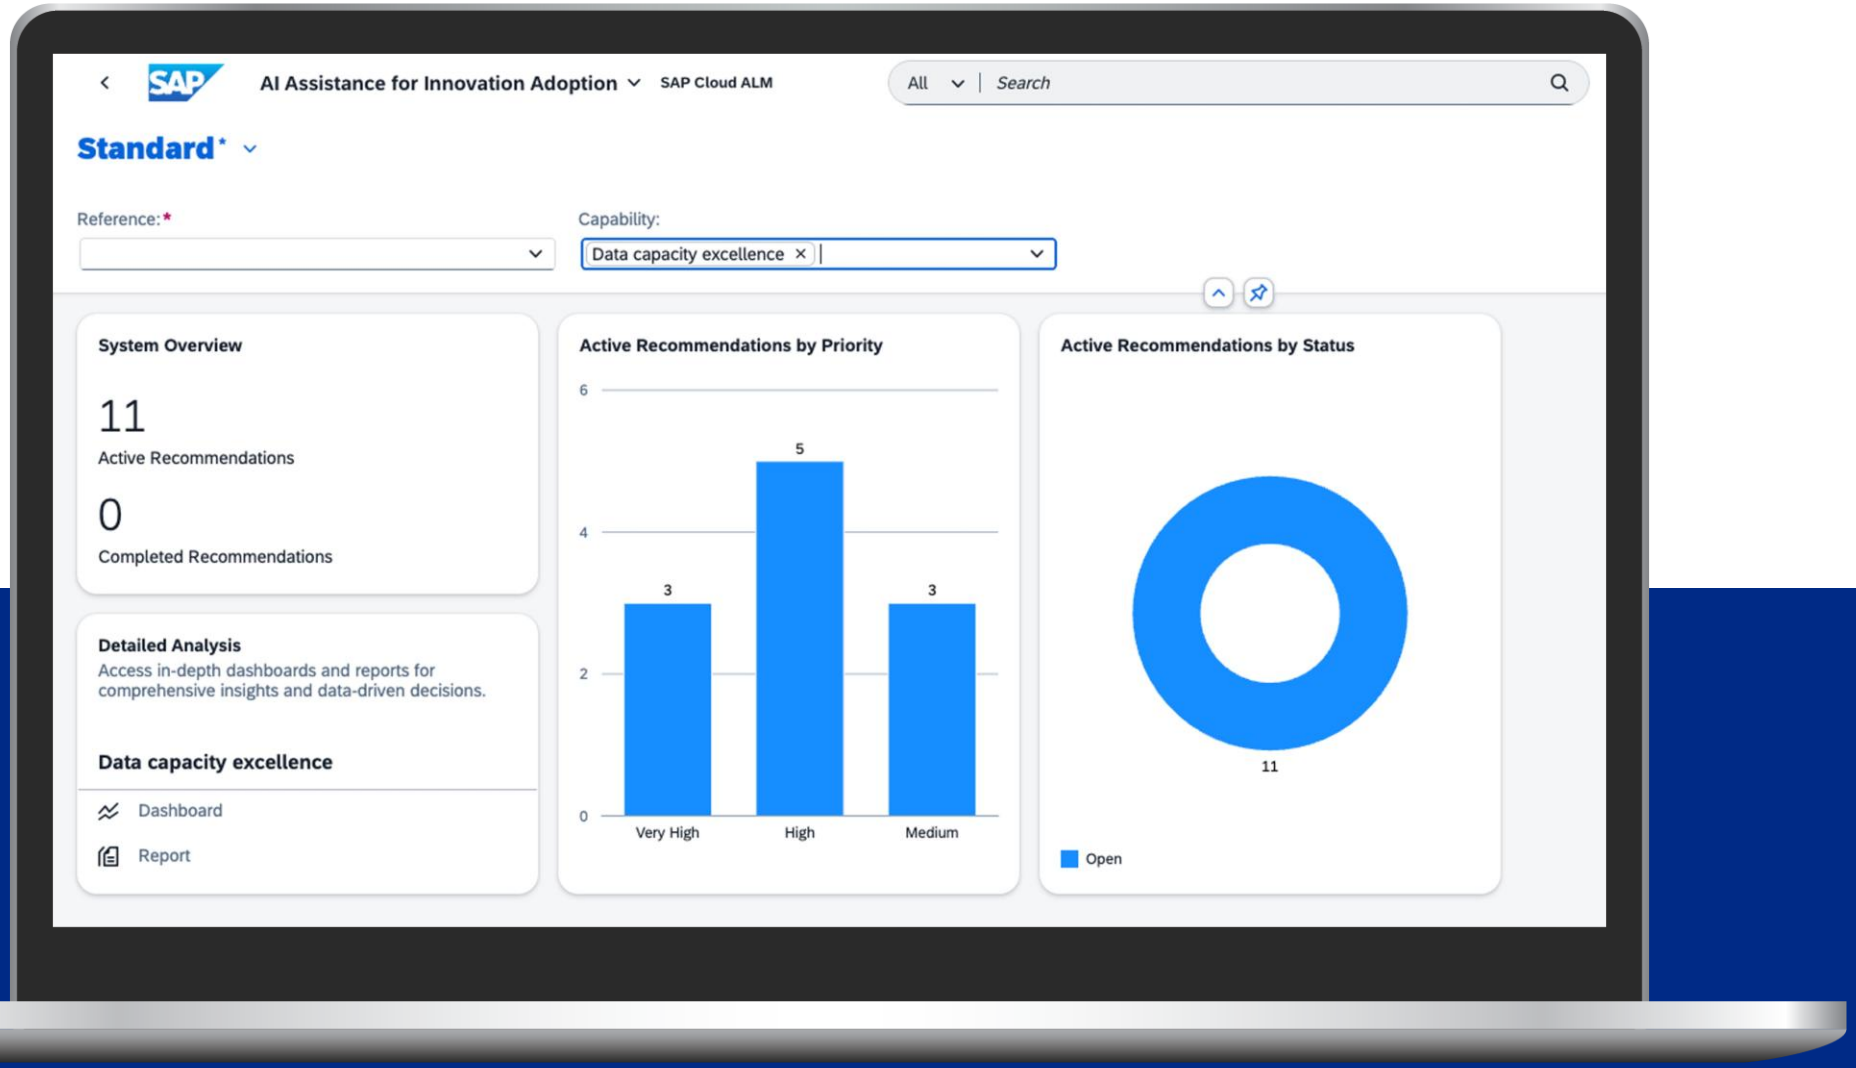

Embedded SAC Dashboard

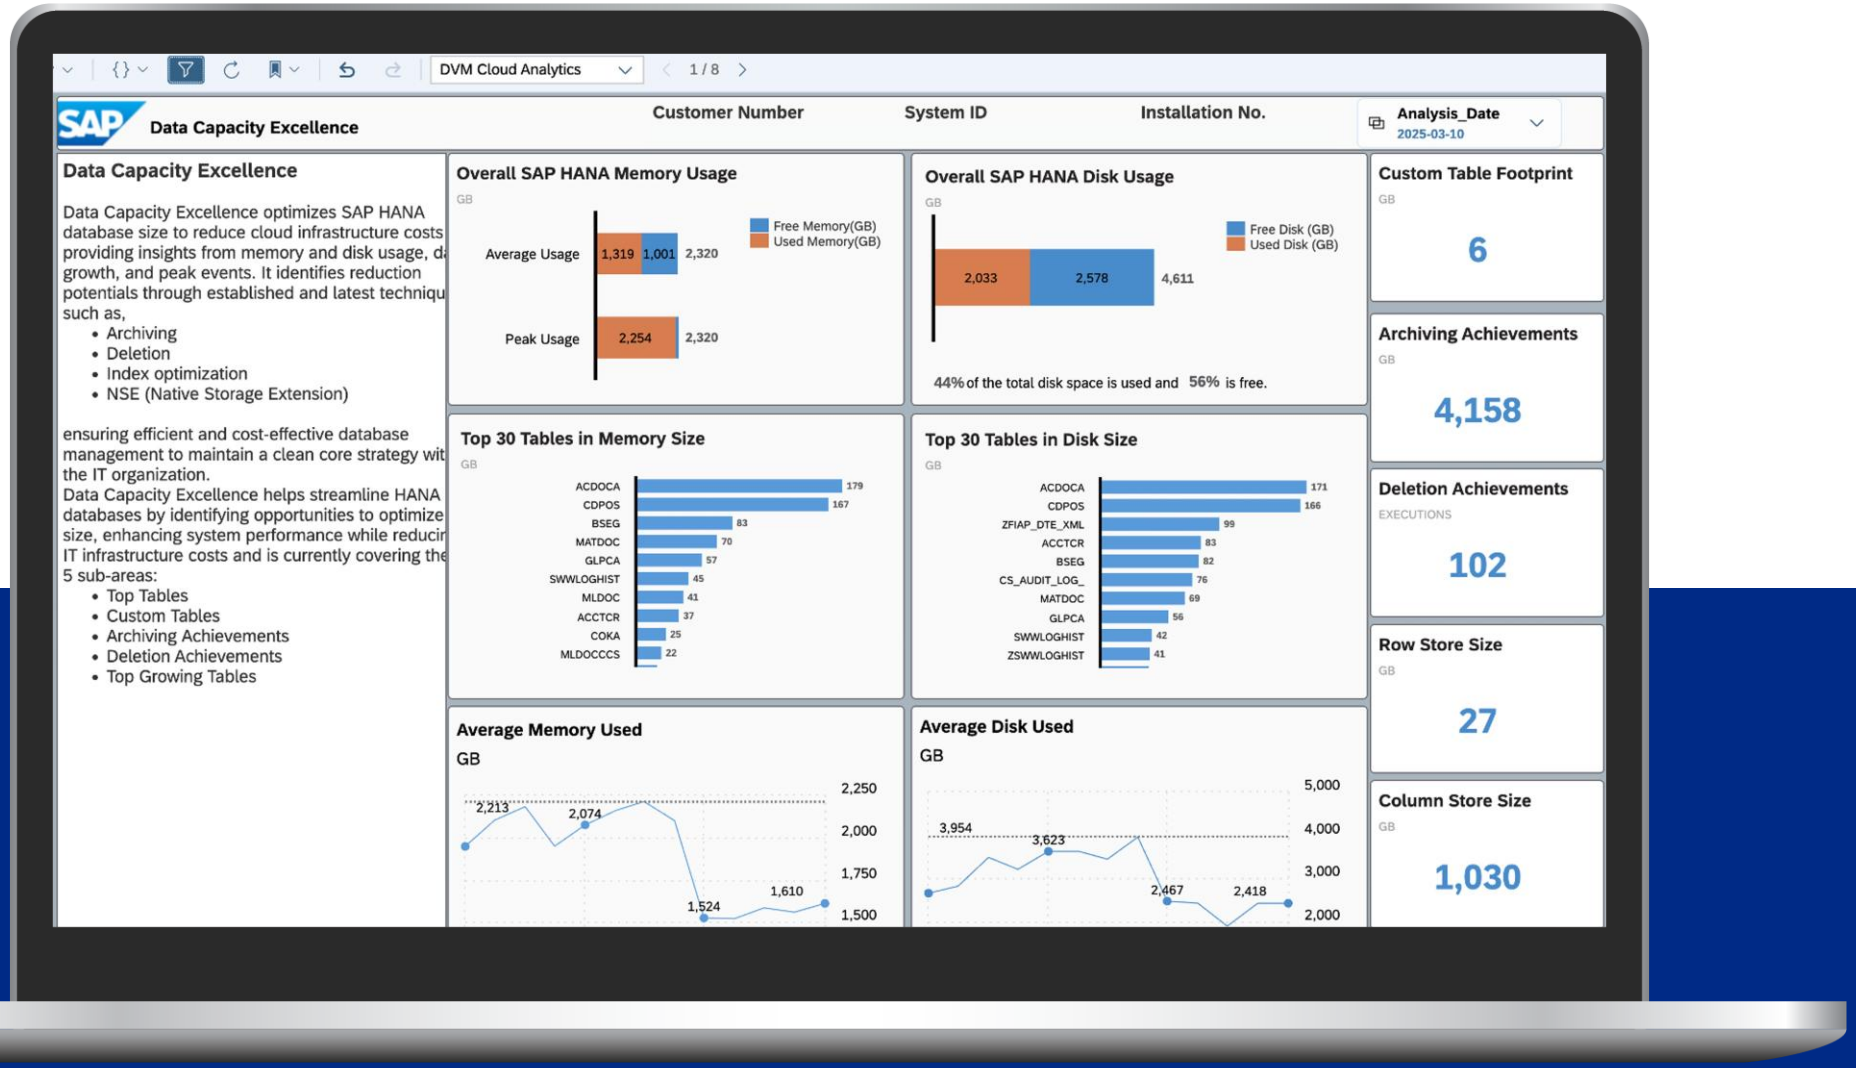

AI assistance for innovation adoption | **Prerequisites**

SAP Cloud ALM

AI assistance for innovation adoption uses SAP Cloud ALM for Service as a collaboration platform. You need a subscription to SAP Cloud ALM and must meet the technical prerequisites for using SAP Cloud ALM for Service

Data Capacity Excellence

- ST-PI 740 version SP31 or higher: Install SAP Note 3703380.
- Assign the background user of the push provider the role `SAP_SDF_ALM_METRIC_PUSH_DVM`.

Operational Excellence of SAP BTP

Only SAP BTP global accounts relating to your customer number are possible to monitor.

Evolution Readiness

- Provide an up-to-date SAP Readiness Check for SAP S/4HANA updates in the relevant SAP Cloud ALM tenant regarding the desired or latest available target release of the production system. See SAP Note 3059197.
- You have provided an up-to-date SAP Readiness Check for SAP Cloud ALM updates in the relevant SAP Cloud ALM tenant regarding the desired or latest available target release of the production system. See SAP Note 3236443.

Enduring Cybersecurity

- ST-PI 740 SP31 or higher: Install SAP Note 3670114.
- Assign the background user of the push provider the role `SAP_SDF_ALM_METRIC_PUSH_CYSEC`.
- Share your system landscape data with SAP using one of the following options: the landscape management database (LMDB) in SAP Solution Manager, SAP Focused Run, or a direct system connection (see SAP Note 2855951).

Technical Efficiency Optimization

- ST-PI 740 SP32 and SAP Note 3676780 / ST-PI 740 SP33 or higher
- ST-A/PI 01W_031 or higher
- SAP S/4HANA single node system
- Assign the background user of the push provider the role `SAP_SDF_ALM_METRIC_PUSH_UTE`.

Success Plans

Subscribe to either the Advanced Success Plan or the Max Success Plan.

Setup | SAP Help Portal

Adoption Potential for RISE with SAP

- `ST-PI 740 SP34`
- Provide an up-to-date SAP Readiness Check for SAP S/4HANA updates in the relevant SAP Cloud ALM tenant regarding the desired or latest available target release of the production system. Custom code analysis is not required. See SAP Note 3059197.
- Activate the SAP EarlyWatch Alert. See SAP Note 1257308.
- Assign the background user of the push provider the role `SAP_SDF_ALM_METRIC_PUSH_ADOPP`.

Extensibility Intelligence

- ST-PI 740 version SP31 or higher
- Implement the relevant SAP Notes based on the release of the ST-PI of your system. See SAP Note 3169647, section 1 for the release-dependent SAP Notes.
- Ensure that usage data from your productive system is available. An 18-months time frame is recommended. See SAP Note 3169647, section 2 for details and referenced SAP Notes.
- Assign the background user of the push provider the role `SAP_SDF_ALM_METRIC_PUSH_ICCM`.

AI assistance for innovation adoption
Data Capacity Excellence

Description

Data Capacity Excellence focuses on keeping your SAP HANA database clean and avoiding unnecessary data. We help you understand what consumes memory and how to reduce it effectively. It is all about smart data volume management analyzing usage patterns, growth trends, indexing, partitioning, and SAP HANA native storage extension.

Business Needs

As a customer you want to …

- Keep your SAP HANA systems clean and optimized
- Avoid paying for unused or unnecessary data
- Improve system performance

Value & Benefits

- Memory savings through optimized data management and efficient storage use
- Key KPIs include SAP HANA memory and disk usage
- Risk reduction by avoiding oversized systems, slow operations, and high costs
- Dashboard offers insights for optimizing SAP HANA memory and system size
- Lower costs and improved performance for SAP HANA customers

Delivery & Approach

Data capacity excellence offers continuous access to an intuitive dashboard that visualizes memory-saving potential across various cleanup methods such as archiving and inverted indexes. Optimization areas are automatically identified with tailored recommendations offered whenever reduction potential is detected.

Content Details

- Top tables
- Custom tables
- Archiving achievements
- Deletion achievements
- Top growing tables
- Inverted individuals
- Archiving potential
- Native Storage Extension (NSE)
- FI Document
- Tables Logging
- Table Partitions
- Deferred Summarizations
- Many more …

PUBLIC

© 2026 SAP SE or an SAP affiliate company. All rights reserved.
See Legal Notice on www.sap.com/legal-notice for use terms, disclaimers, disclosures, or restrictions related to this material.

Data Capacity Excellence

Name	Description
Document Types (Reduction Potential)	DVM – Check the Reduction Potential of the Document Types of an S/4 system
Document Types (Memory & Disk Growth)	DVM – Check the Memory & Disk Growth of the document types of an S/4 system
Tables (Usage)	DVM – Check the read count and the NSE configurations of the top tables of an S/4 system.
Tables (Partitioning)	DVM – Check current partitioning of tables in an S/4 System
Tables (Indexes)	DVM – Check whether indexes are required or redundant for any of the top tables of an S/4 system.
HANA Disk	DVM – Check the HANA Disk of an S/4 system
HANA Memory	DVM – Check the HANA Memory of an S/4 system

Key Techniques in Data Volume Management

Application-Layer

Data Avoidance
Deactivating the creation of data in the first place

Data Deletion
Deleting records which are no longer required by business processes or legal compliance

Data Summarization
Aggregating data at higher level if enough for the business

Data Archiving
Remove data from the database and store it in an external server with limited access

Database-Layer

Native Storage Extension (NSE)
Moving data from the in-memory part of a SAP HANA database to the disk part

Index Optimization
Remove redundant indexes of a table

Partitioning
Split and distribute tables with different load priority

Issue and Actions Management in Cloud ALM

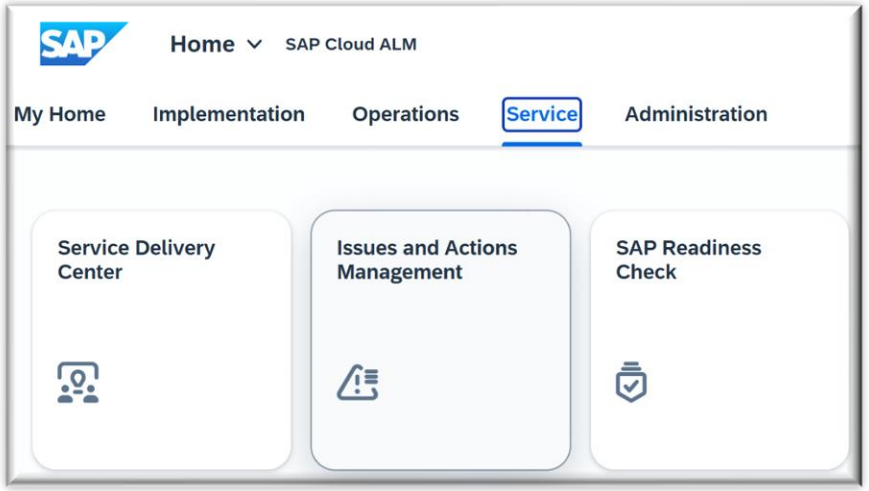

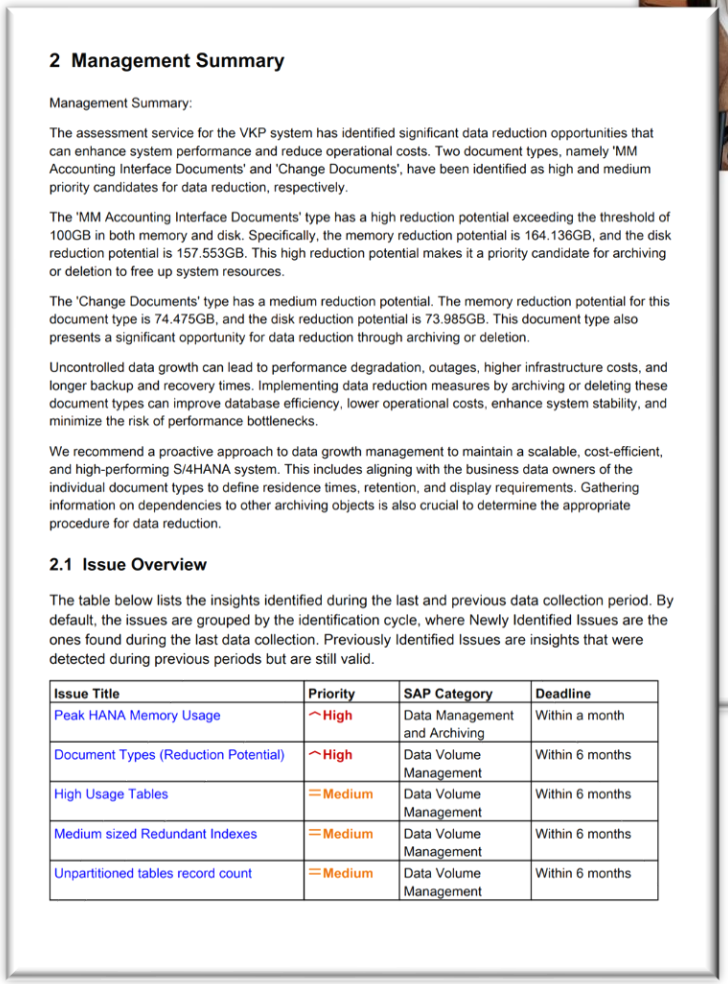

2 Management Summary

Management Summary:

The assessment service for the VKP system has identified significant data reduction opportunities that can enhance system performance and reduce operational costs. Two document types, namely 'MM Accounting Interface Documents' and 'Change Documents', have been identified as high and medium priority candidates for data reduction, respectively.

The 'MM Accounting Interface Documents' type has a high reduction potential exceeding the threshold of 100GB in both memory and disk. Specifically, the memory reduction potential is 164.136GB, and the disk reduction potential is 157.553GB. This high reduction potential makes it a priority candidate for archiving or deletion to free up system resources.

The 'Change Documents' type has a medium reduction potential. The memory reduction potential for this document type is 74.475GB, and the disk reduction potential is 73.985GB. This document type also presents a significant opportunity for data reduction through archiving or deletion.

Uncontrolled data growth can lead to performance degradation, outages, higher infrastructure costs, and longer backup and recovery times. Implementing data reduction measures by archiving or deleting these document types can improve database efficiency, lower operational costs, enhance system stability, and minimize the risk of performance bottlenecks.

We recommend a proactive approach to data growth management to maintain a scalable, cost-efficient, and high-performing S/4HANA system. This includes aligning with the business data owners of the individual document types to define residence times, retention, and display requirements. Gathering information on dependencies to other archiving objects is also crucial to determine the appropriate procedure for data reduction.

2.1 Issue Overview

The table below lists the insights identified during the last and previous data collection period. By default, the issues are grouped by the identification cycle, where Newly Identified Issues are the ones found during the last data collection. Previously Identified Issues are insights that were detected during previous periods but are still valid.

Issue Title	Priority	SAP Category	Deadline
Peak HANA Memory Usage	High	Data Management and Archiving	Within a month
Document Types (Reduction Potential)	High	Data Volume Management	Within 6 months
High Usage Tables	Medium	Data Volume Management	Within 6 months
Medium sized Redundant Indexes	Medium	Data Volume Management	Within 6 months
Unpartitioned tables record count	Medium	Data Volume Management	Within 6 months

Service Name	SAP Data Capacity Excellence
Customer Name	
System	VKP
Date of Report	Mar 19, 2025

Issue and Actions Management in Cloud ALM

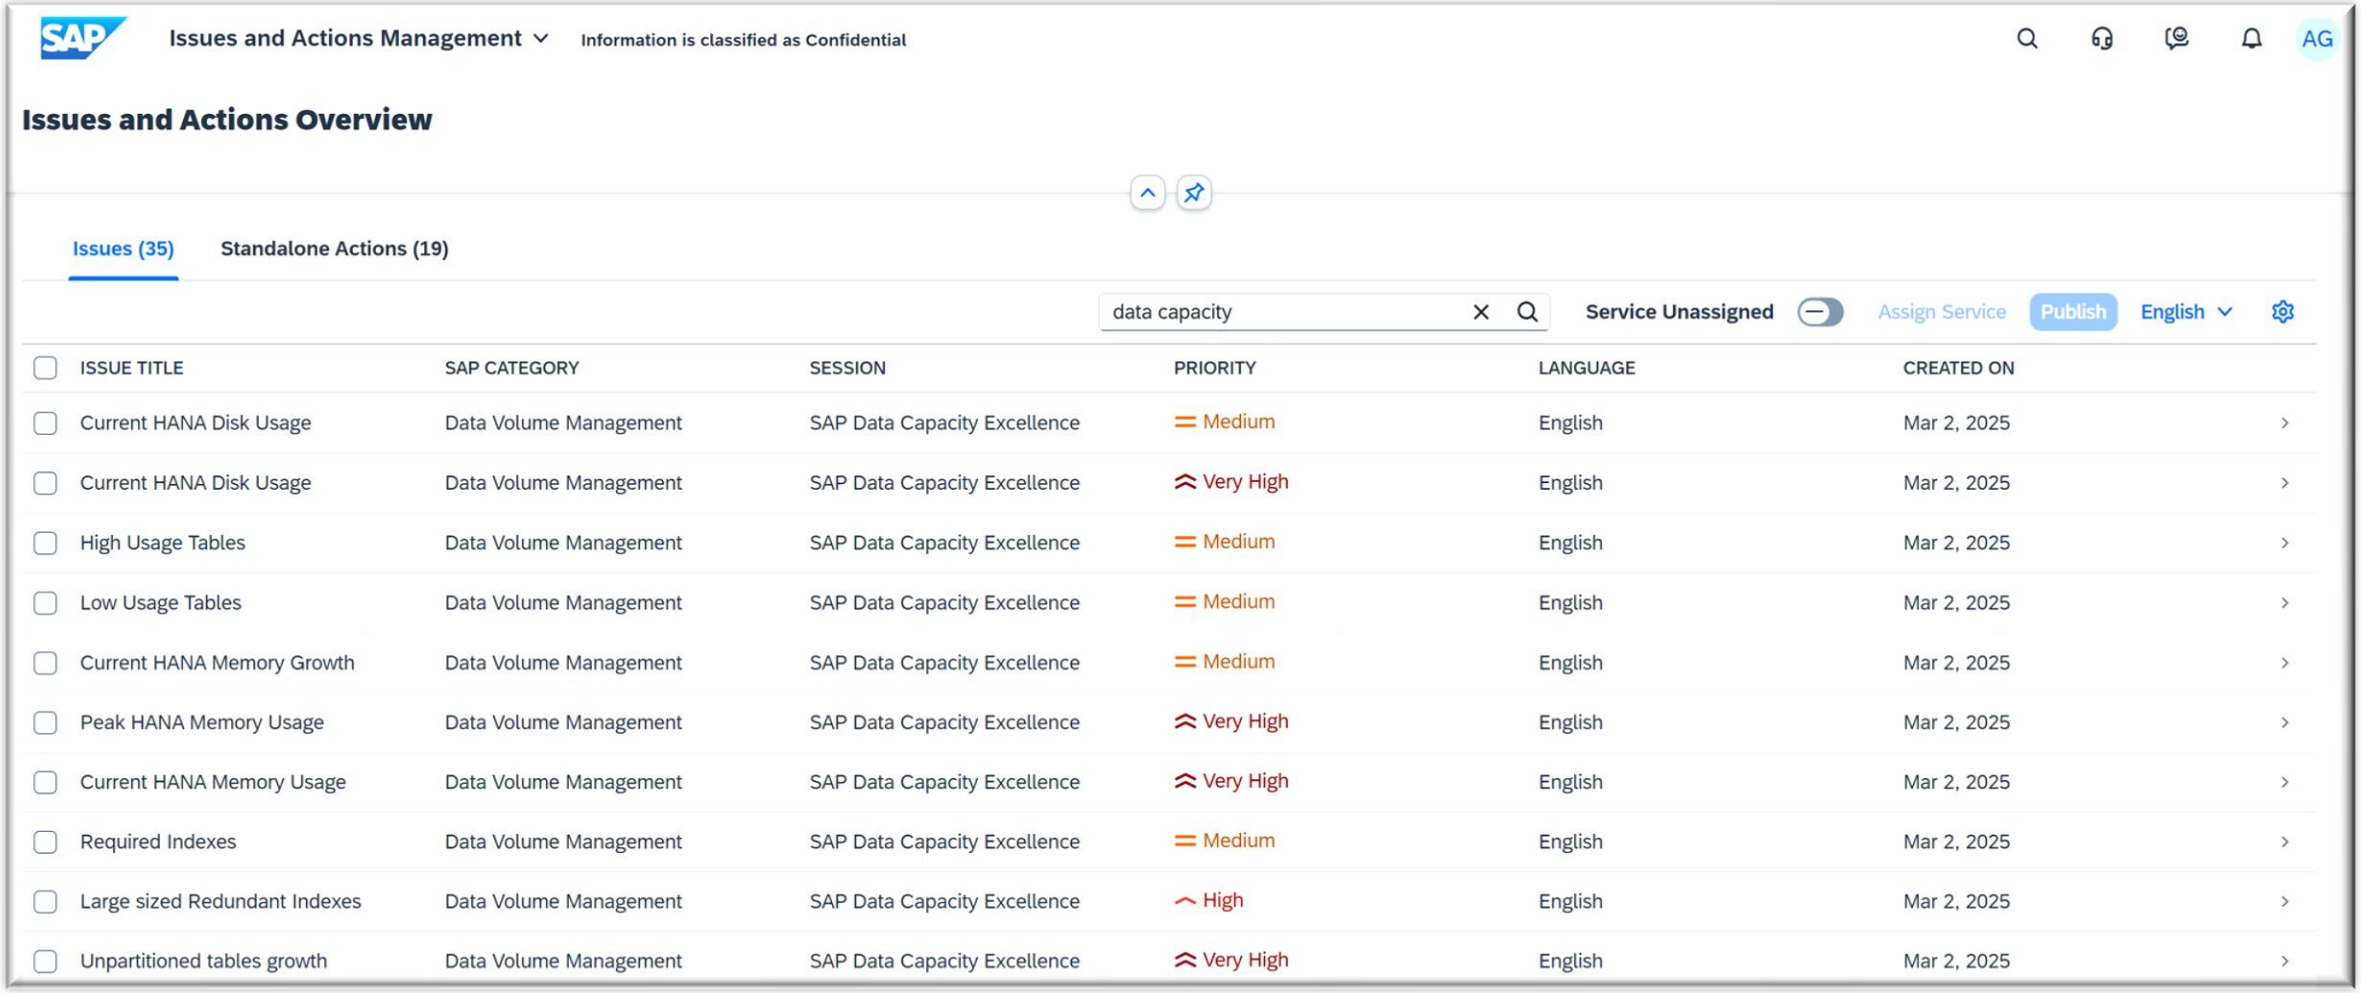

ISSUE TITLE	SAP CATEGORY	SESSION	PRIORITY	LANGUAGE	CREATED ON
Current HANA Disk Usage	Data Volume Management	SAP Data Capacity Excellence	Medium	English	Mar 2, 2025
Current HANA Disk Usage	Data Volume Management	SAP Data Capacity Excellence	Very High	English	Mar 2, 2025
High Usage Tables	Data Volume Management	SAP Data Capacity Excellence	Medium	English	Mar 2, 2025
Low Usage Tables	Data Volume Management	SAP Data Capacity Excellence	Medium	English	Mar 2, 2025
Current HANA Memory Growth	Data Volume Management	SAP Data Capacity Excellence	Medium	English	Mar 2, 2025
Peak HANA Memory Usage	Data Volume Management	SAP Data Capacity Excellence	Very High	English	Mar 2, 2025
Current HANA Memory Usage	Data Volume Management	SAP Data Capacity Excellence	Very High	English	Mar 2, 2025
Required Indexes	Data Volume Management	SAP Data Capacity Excellence	Medium	English	Mar 2, 2025
Large sized Redundant Indexes	Data Volume Management	SAP Data Capacity Excellence	High	English	Mar 2, 2025
Unpartitioned tables growth	Data Volume Management	SAP Data Capacity Excellence	Very High	English	Mar 2, 2025

Example of the issue “Unpartitioned tables growth”

Issues and Actions Management Information is classified as Confidential AG

Issues and Actions Management / Issue

Unpartitioned tables growth

Edit Delete English (Original)

Go Live Date: Jan 28, 2025

Issue Status: Draft

General Info Recommended Actions Attachments Change Log

This check evaluates record growth over the past 42 days to identify tables that, if the current growth trend continues, are likely to reach critical thresholds. In System A17 we have found 2 tables which are unpartitioned and which will reach the critical 2 billion record count limit in less than 3 months based on their current record count growth.

These are the affected tables including their current record count:

Table Name	Record Count
ACCTCR	1,622,889,115
BSEG	1,372,706,830

Example of the issue “Unpartitioned tables growth”

Issues and Actions Management Information is classified as Confidential AG

Issues and Actions Management / Issue

Unpartitioned tables growth

Edit Delete English (Original)

General Info Recommended Actions Attachments Change Log

Recommended Actions

Actions (1)

Action Title
Partition tables with high record count

Priority:
Very High

Due Date:
–

Assigned Role:
System Administrator

Deadline:
Within a month

Description:

Implement partitioning for the Identified tables.

1. **Number of required partitions:** Identify the number of required partitions based on the current and expected size of the table. The size of a partition should be around 15-30 GB. A good rule of thumb is 250.000.000 – 500.000.000 records per partition.
2. **Partition type:** Choose between
   - HASH: Simply divides the records by half, by third or any number specified in the number of partitions. Easy to implement.
   - RANGE: Partitions and its ranges are exactly defined. Allows the separation of old and new data. Harder to implement.
3. **Partition key:** Choose the table column by which the partitioning should be created. This column should be a highly used primary key by which the table is accessed.
   - **HASH:** The column should have a lot of different values. (e.g., CHANGENR or BELNR)
   - **RANGE:** The column should be incremental (e.g., date columns like GJAHR or CHANGENR)
4. Execute the partitioning command in HANA DB. **Important:** The execution of (re-)partitioning for a table requires extra preparation and attention. This is because the execution requires extra resources, and the affected tables are at certain points in time unavailable. Please consider Go-Live Support from SAP Premium Engagement.
5. For more information see also 2044468 - FAQ: SAP HANA Partitioning.

Example of the issue “Unpartitioned tables growth”

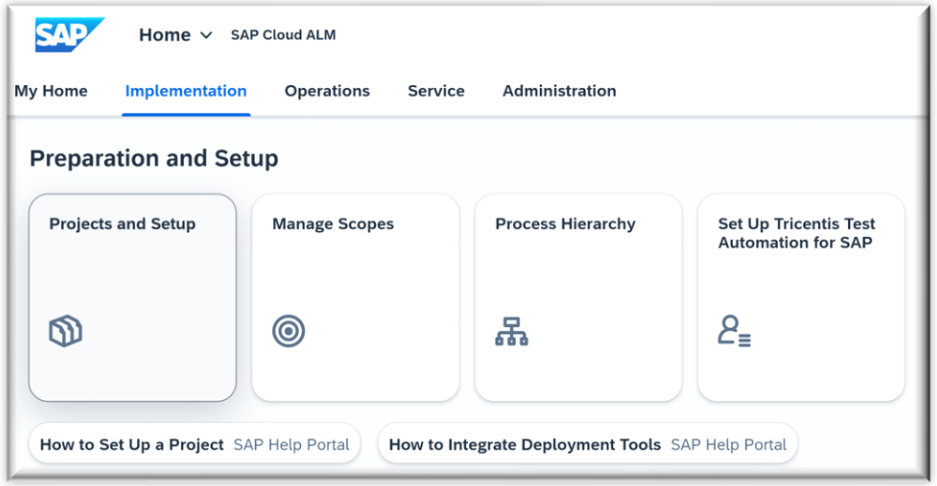

If the recommended Action description of the identified issues is not sufficient enough, then **Project and Setup** App in CALM can support the implementation with detailed steps and additional knowledge on implementation guidance.

Example of the issue “Unpartitioned tables growth”

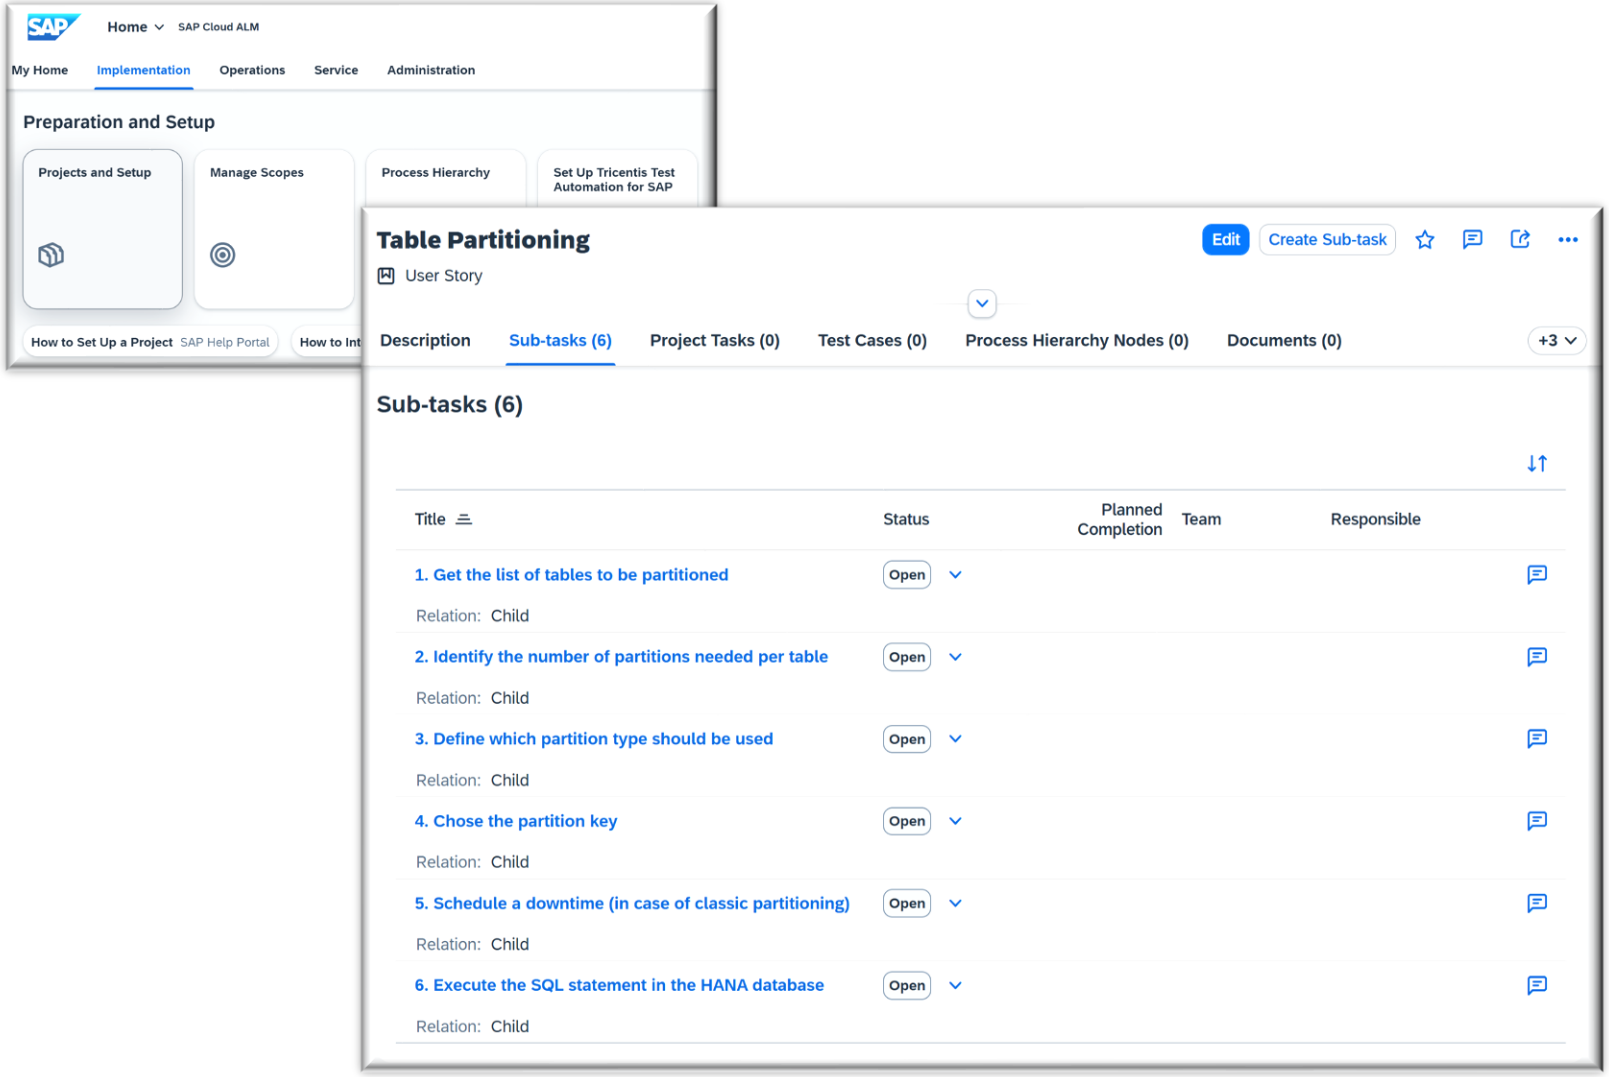

Example of the issue “Unpartitioned tables growth”

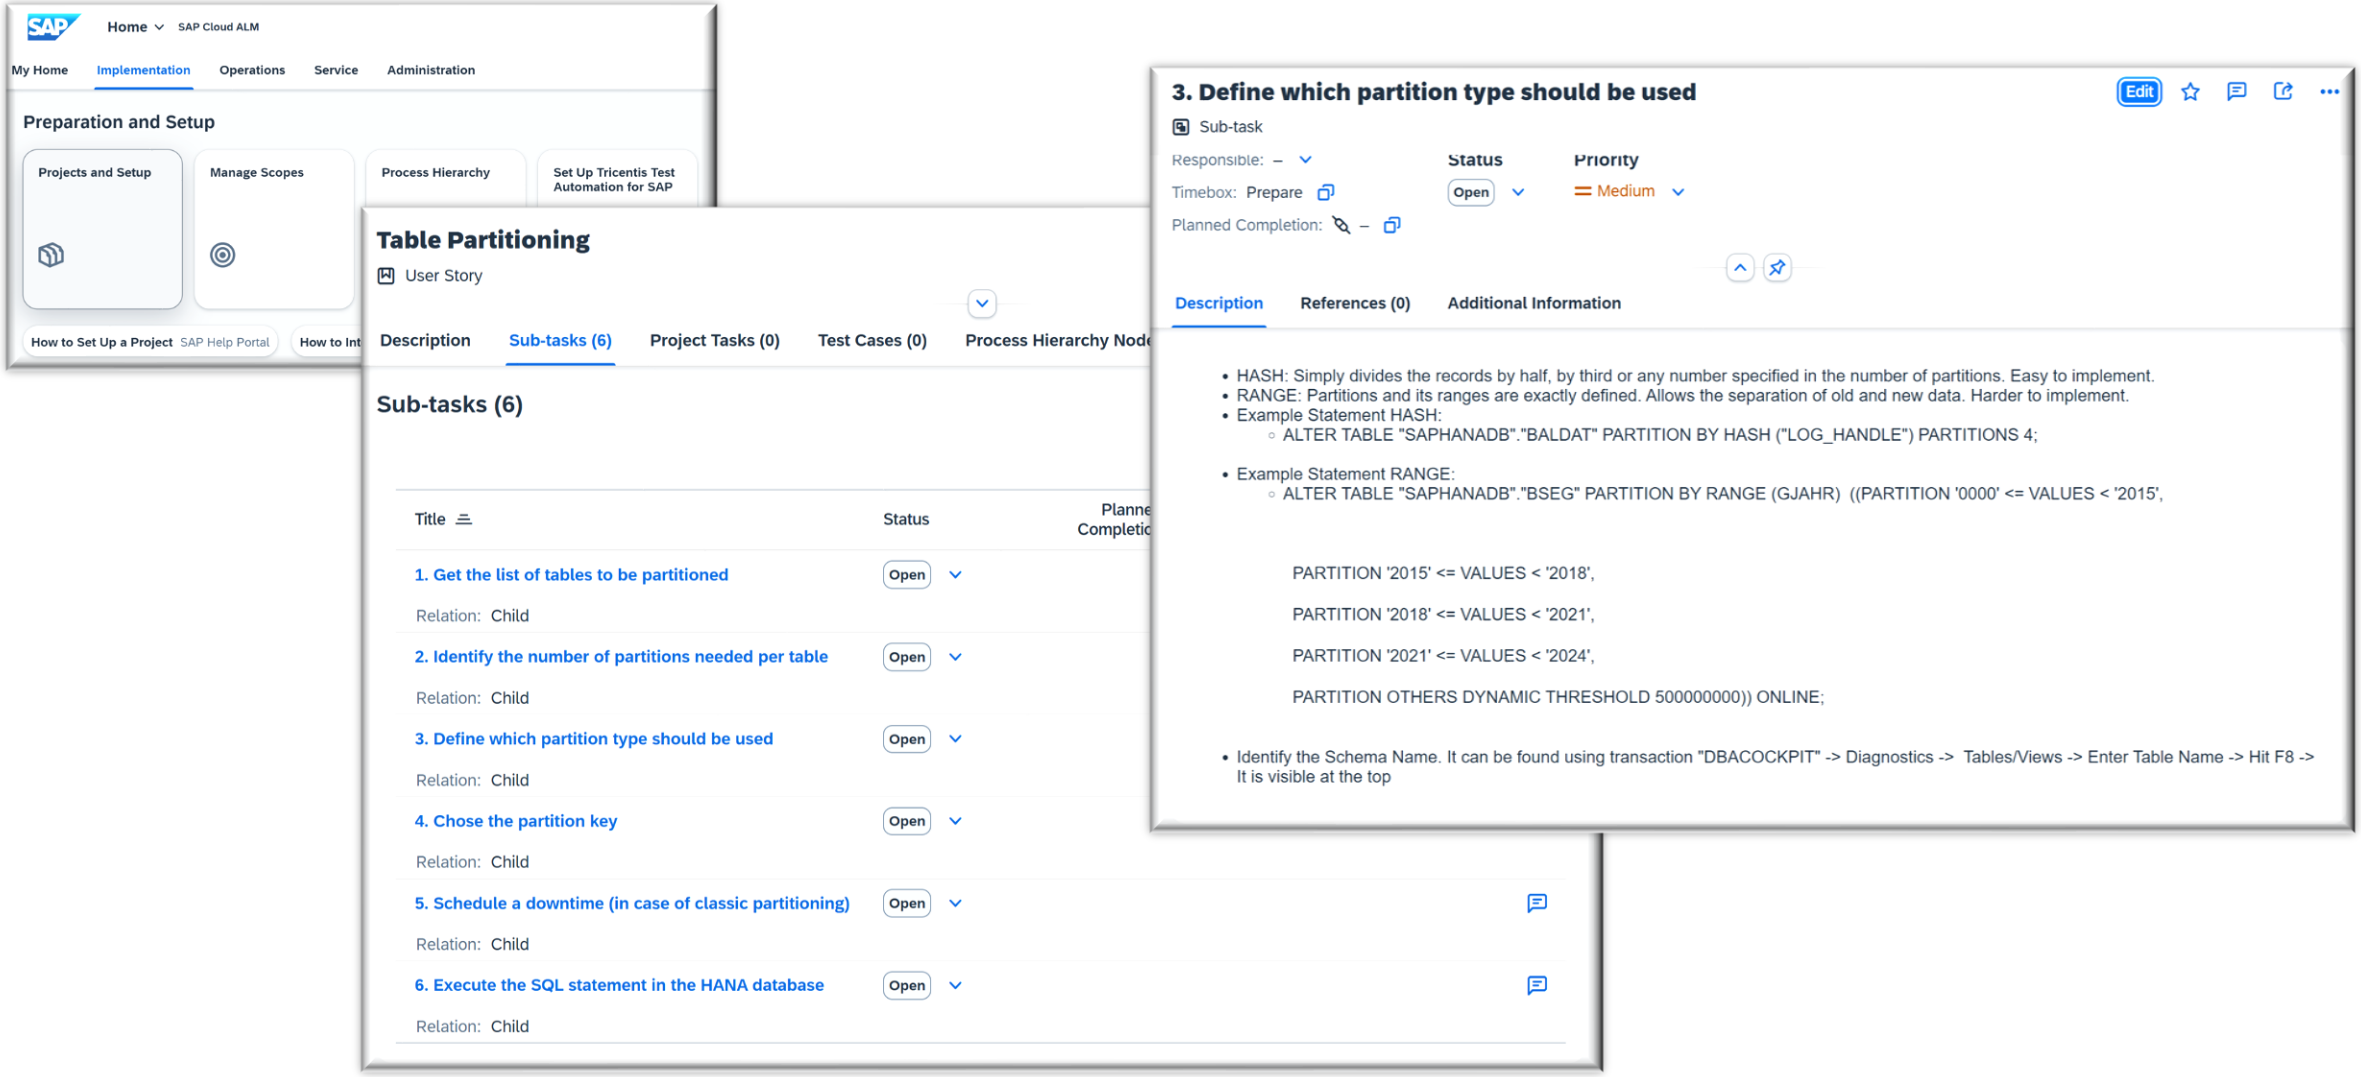

Example of the issue “Document Types (Reduction Potential)”

Issues and Actions Management / Issue

Document Types(Reduction Potential)

Edit Delete English (Original)

General Info | Recommended Actions | Attachments | Change Log

Priority:
High

Assigned Role:
System Administrator

Deadline:
Within 6 months

SAP Category:
Data Volume Management

Due Date:
–

System:
–

Description:

This check analyses the document types in your S/4 Hana system to identify reduction potential in both memory and disk storage. The goals is to determine whether archiving or deletion should be prioritized to optimize system performance and prevent unnecessary data growth.
The reduction potential is calculated by analysing historical data distributions in your system and applying SAP Best Practice residence times. These residence times define the period after which data is typically no longer required for daily operations and can be archived or deleted to prevent continuous database growth. Without proactive data management, the database size would steadily increase, leading to higher infrastructure costs and potential performance issues

.

Doc Type	Memory Reduction Potential (GB)
MM Accounting Interface Documents	164.13
Change Documents	74.47

Doc Type	Disk Reduction Potential (GB)
MM Accounting Interface Documents	157.55
Change Documents	73.98

Example of the issue “Document Types (Reduction Potential)”

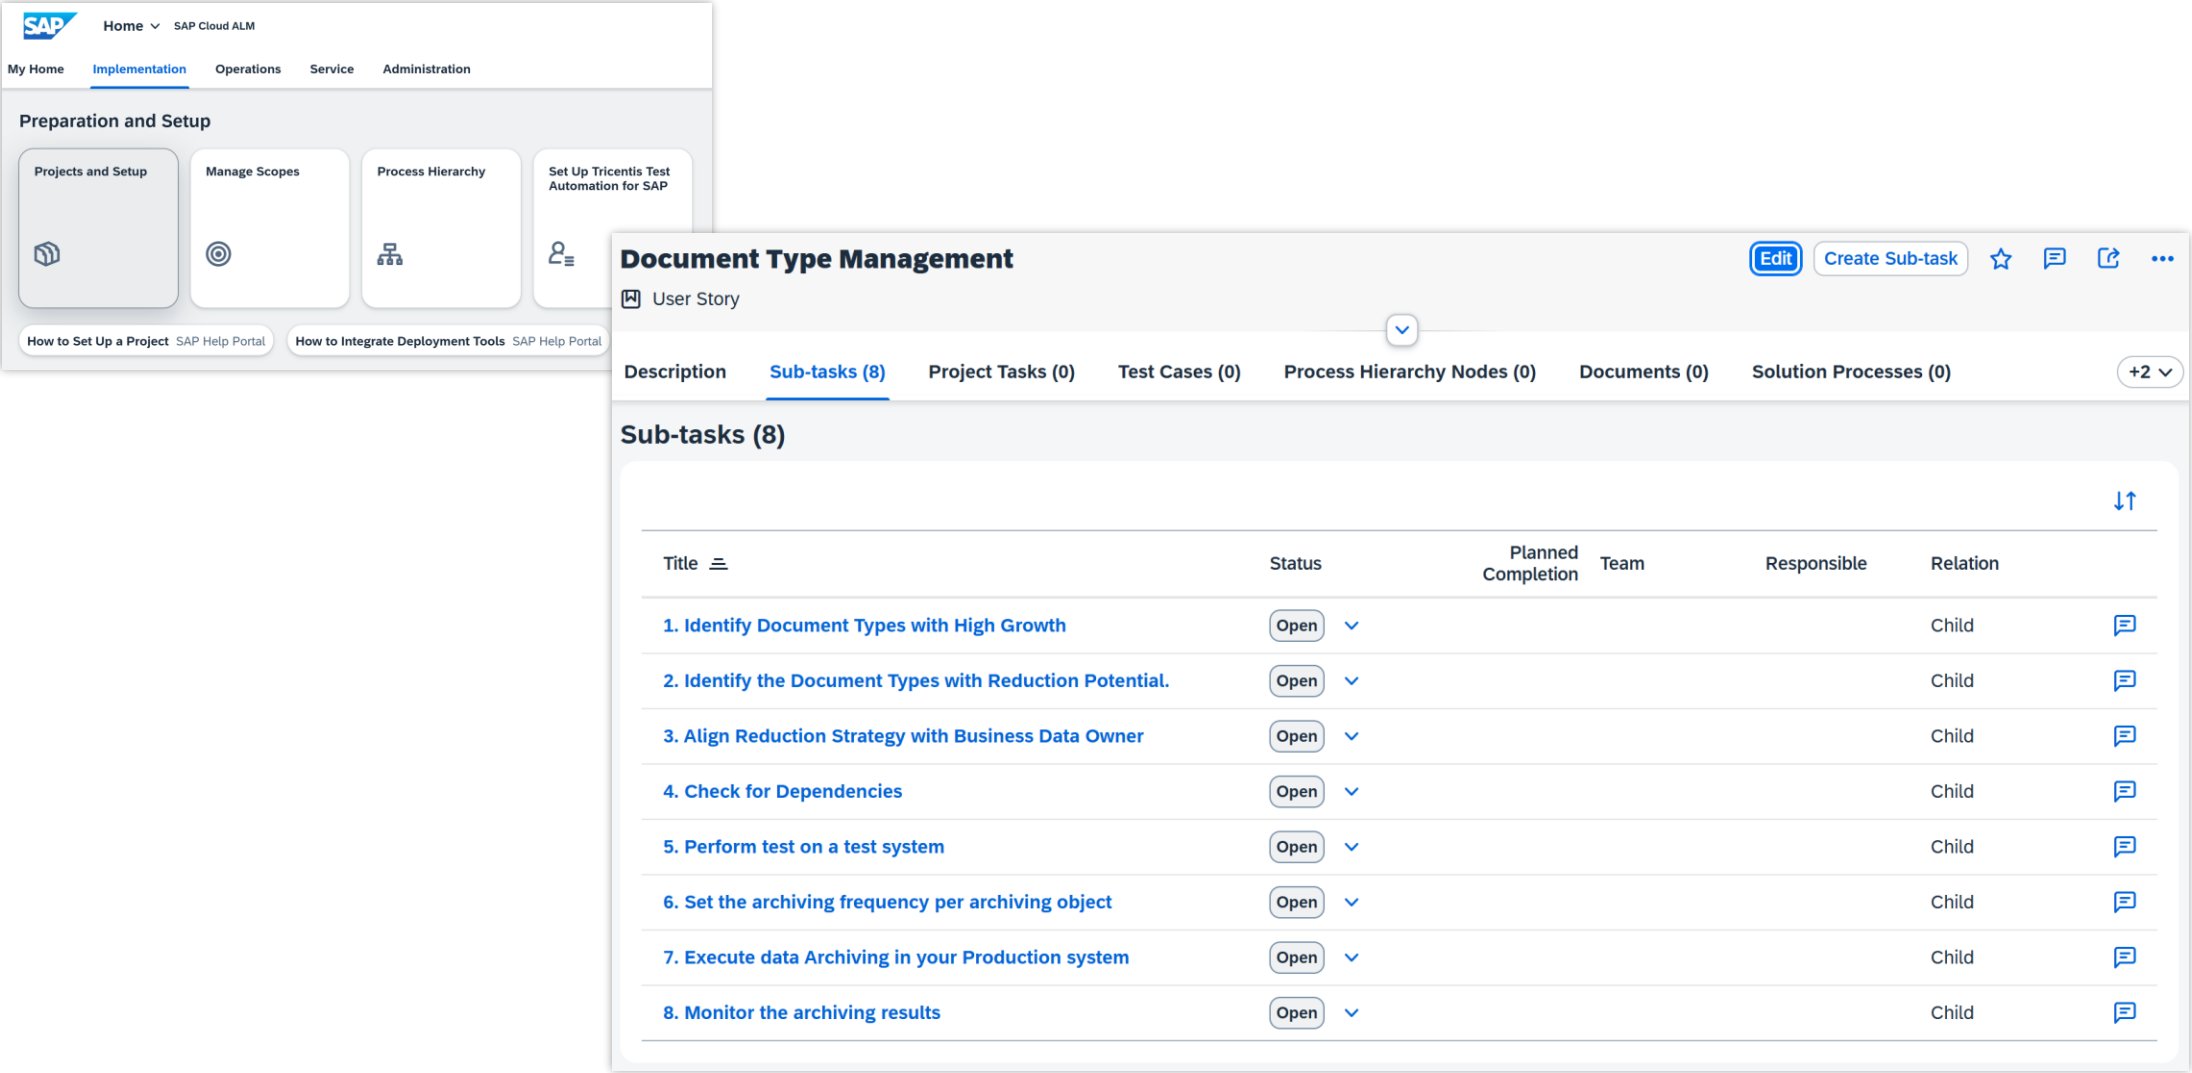

Example of the issue “Document Types (Reduction Potential)”

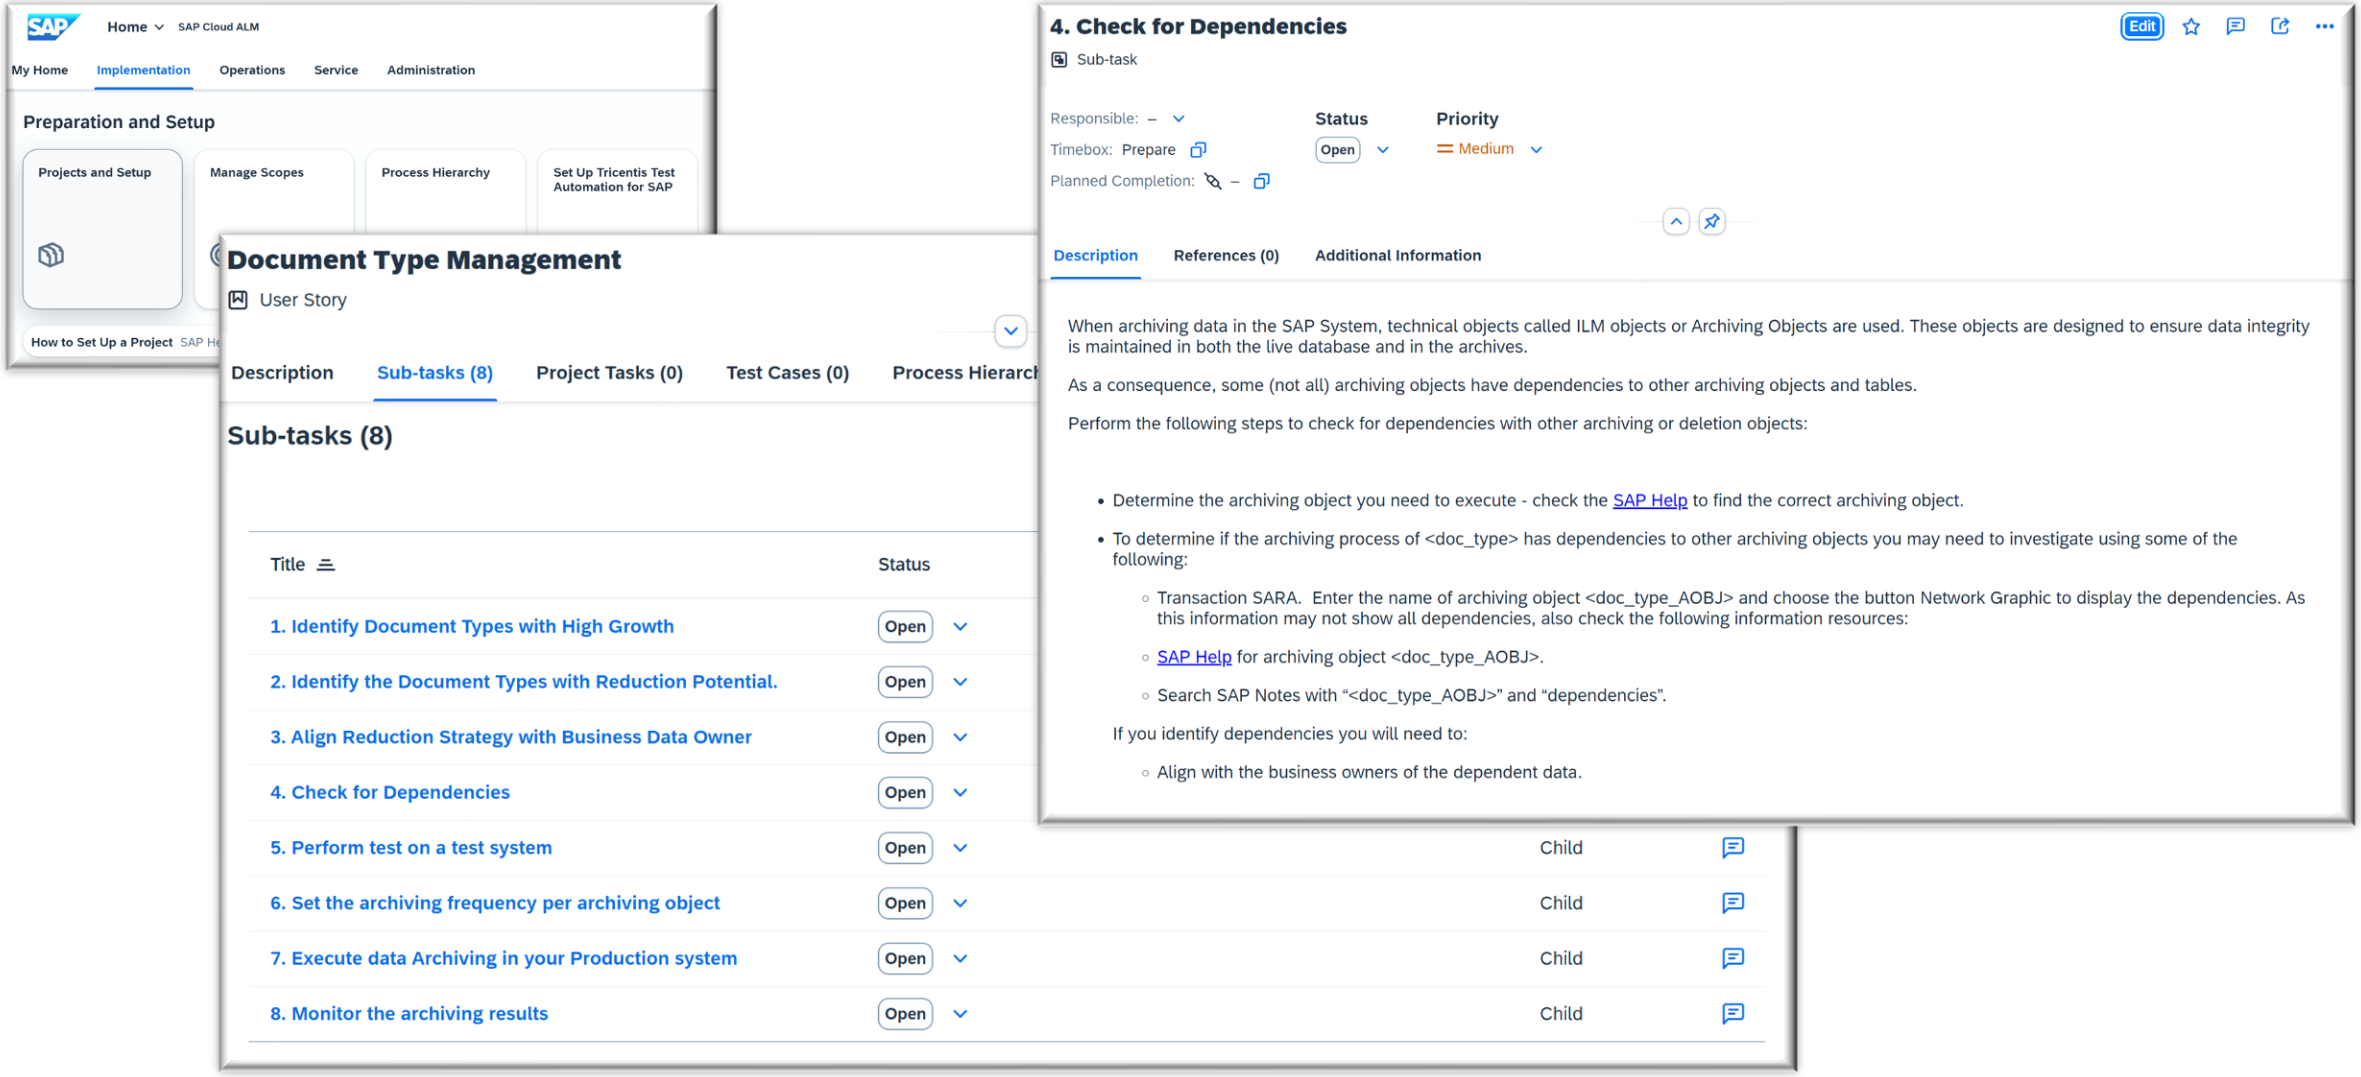

Example of the issue “Large sized Redundant Indexes”

Issues and Actions Management Information is classified as Confidential AG

Issues and Actions Management / Issue

Large sized Redundant Indexes

Edit Delete English (Original)

General Info Recommended Actions Attachments Change Log

Priority:
High

Assigned Role:
System Administrator

Deadline:
Within 6 months

SAP Category:
Data Volume Management

Due Date:
–

System:
–

Description:

This check identifies redundant indexes in your S/4 system based on the Inverted Individual Costs (INDIV_COSTS) of the tables. An index is considered redundant if the INDIV_COSTS of at least one column in the table are below 200 and an index exists for that table.
In system A17, we have found 2 tables containing large-sized redundant indexes exceeding 50 GB in disk or memory usage. These indexes are consuming significant resources and can be optimized. The following tables are affected, including their potential reduction in Memory (in GB):
Please see SAP Note 2600076 for more information.

Table Name	Memory Reduction Potential (GB)
EDID4	64.4
CDPOS	51.0

Business Impact:

Reduce HANA Disk and Memory Usage by removing the redundant indexes.

Integrated SAC Dashboard for Transparency & Monitoring

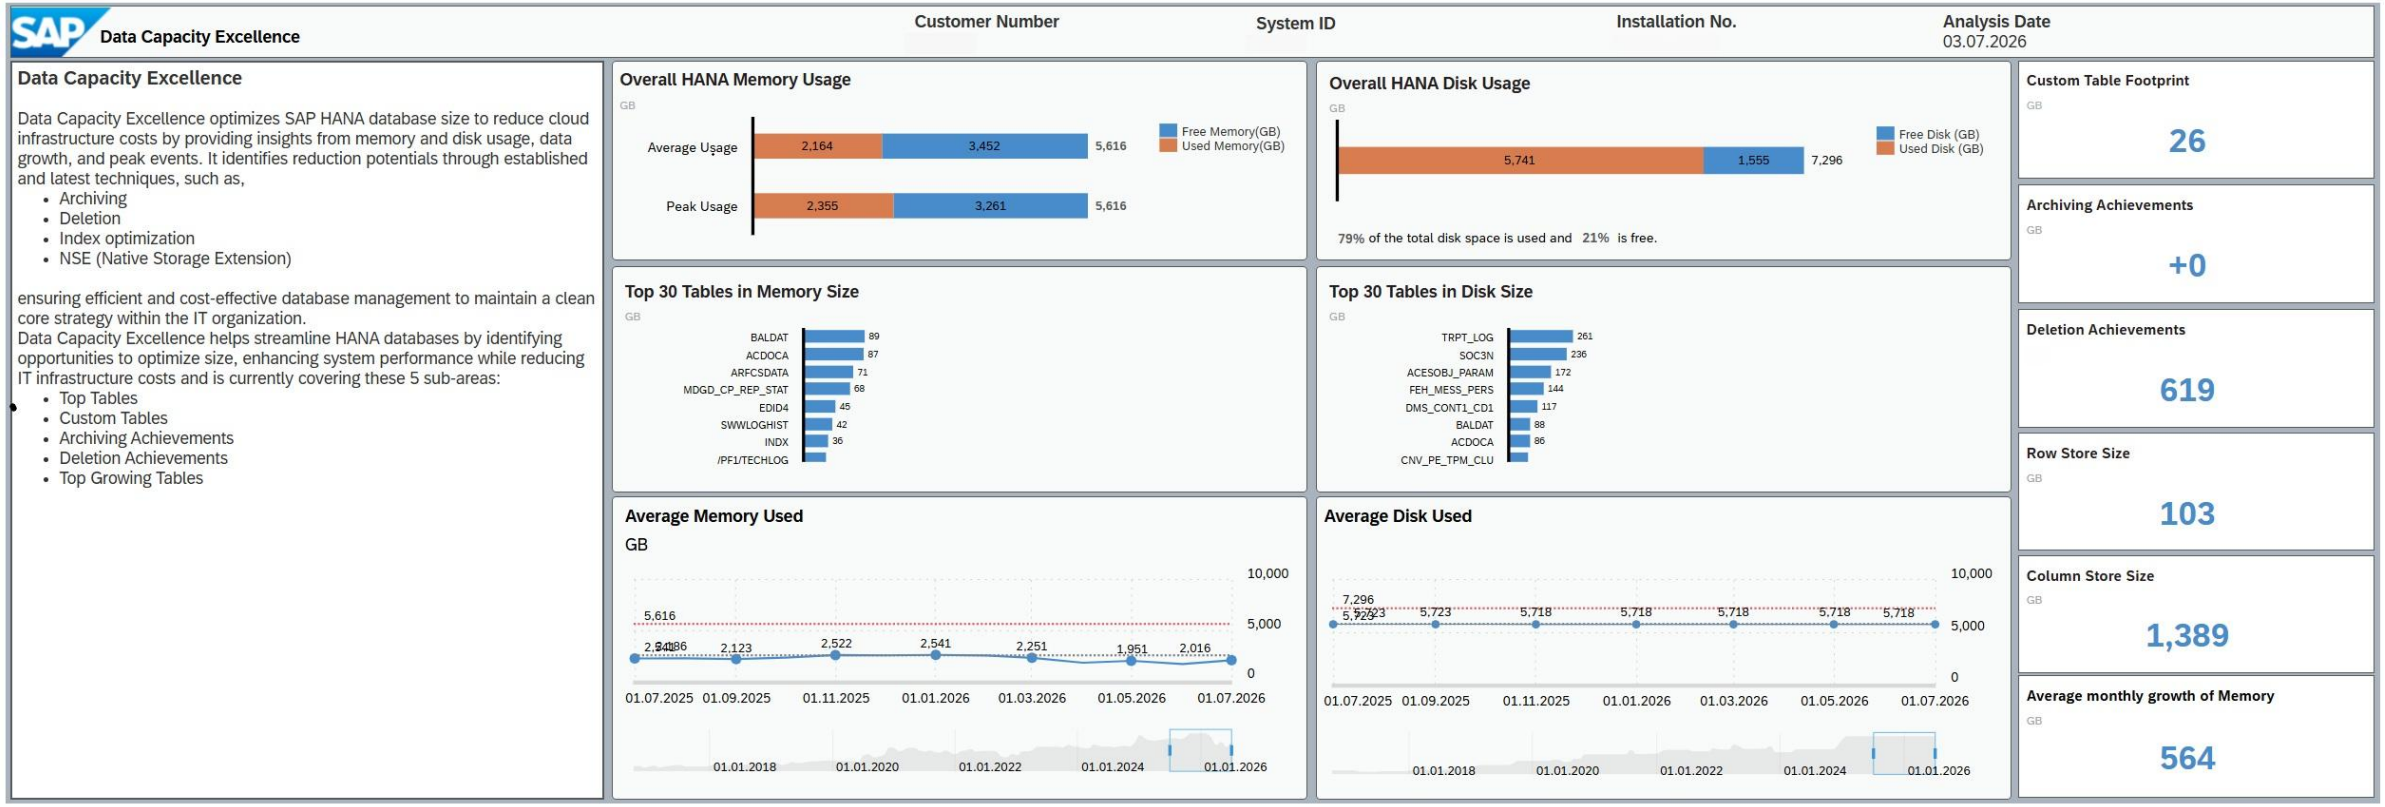

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Top Tables

This page provides a detailed overview of the largest tables in the SAP S/4HANA system, offering insights into the system's memory and disk usage. The following key measures are displayed for analysis:

- **MEM_GB**: The memory footprint of each table, measured in gigabytes.
- **DISK_GB**: The size of the table on disk, measured in gigabytes.
- **RECORD_COUNT**: The total record count for each table.
- **PART_COUNT**: The number of partitions for a table. A value of "0" indicates that the table is not partitioned.
- **LOB_DISK_GB**: The size of large object (LOB) data stored on disk.
- **LOB_MEM_GB**: The size of large object (LOB) data stored in memory.

The **TABLE_NAME** filter allows you to search for specific tables to focus your analysis.

This page displays the top 1,000 tables ranked by memory footprint, offering detailed visibility into high-impact data objects and their associated memory and storage resource utilization within the system.

TABLE_NAME

- All
- $TA_GML_INDEX_MATGRP_MO2
- $TA_GML_INDEX_MATGRP_REB
- $TA_INDEX_MATGRP_REB
- $TM_DOCUMENTS_GML_INDEX_MATGRP_REB_TEST
- $TM_MATRIX_GML_INDEX_MATGRP_MO2

Apply Selections

Top Memory consuming Tables

MEM_GB 1

TABLE_NAME	DESCRIPTION	MEM_GB	DISK_GB	RECORD_COUNT	PART_COUNT	LOB_DISK_GB	LOB_MEM_GB
BALDAT	Application Log: Log data	88.67	88.40	229,460,388.00	1.00	0.00	0.00
ACDOCA	Universal Journal Entry Line Items	87.01	86.05	2,698,831,148.00	4.00	0.00	0.00
ARFCSDATA	ARFC Call Data (Callers)	71.46	78.53	40,678,844.00	1.00	0.00	0.00
MDGD_CP_REP_STAT	Change Pointer Replication status per target system	67.95	67.13	1,636,163,636.00	1.00	0.00	0.00
EDID4	IDoc Data Records from 4.0 onwards	45.29	44.96	1,430,468,840.00	1.00	0.00	0.00
SWWLOGHIST	Workflow Runtime: History of a Work Item	42.05	41.69	722,309,050.00	1.00	0.00	0.00
INDX	System Table INDX	35.81	36.40	21,417,306.00	1.00	0.00	0.00
/PF1/TECHLOG	Management table for technical logging	31.86	31.54	440,283,663.00	1.00	0.00	0.00
MOFF	Outstanding Article Master Records	29.61	29.23	788,752,501.00	1.00	0.00	0.00
SWWCNTP0	Workflow Container: XML Database (P0)	26.79	26.85	72,236,230.00	1.00	0.00	0.00
FMUSFGA	Actual line item table for US Federal Government.	22.74	39.95	1,019,399,003.00	1.00	0.00	0.00
CDPOS	Change document items	22.46	22.53	306,592,321.00	2.00	0.00	0.00
SWFREVTCNT	Container: XML Runtime Data	22.17	22.31	96,039,931.00	1.00	0.00	0.00
AUSP	Characteristic Values	21.81	21.58	467,349,848.00	1.00	0.00	0.00
/SCA/KT1SPM	Key figure history for time series data 1SPM	20.66	50.03	935,222,571.00	1.00	0.00	0.00
/PF1/DB_OBJ_REL	PE Business Object Relations	20.15	19.94	569,184,376.00	1.00	0.00	0.00
/SMB/SBA1_T_FILE	SBA1-PME table	17.70	17.45	499,463,193.00	1.00	0.00	0.00
WBCROSSGT	Index for Global Types - Where-Used List Workbench	15.84	15.54	246,943,664.00	1.00	0.00	0.00
FQM_FLOW	FQM Flows	15.20	15.06	208,040,892.00	1.00	0.00	0.00
ZDMLT_RENUM_FMUS	USFG - Renumbered DOCLN	15.00	22.55	909,867,000.00	1.00	0.00	0.00
DBTABLOG	Log Records of Table Changes	13.67	36.94	196,898,775.00	1.00	23.31	0.00
JEST	Individual Object Status	13.40	13.10	464,525,656.00	1.00	0.00	0.00
/SAPAPO/DRPMCHNG	Manual Changes to Scheduling Agreements / Purchase Orders	13.27	36.67	836,605,450.00	1.00	0.00	0.00
MDGD_CP	MDG Change Pointer	13.27	13.00	466,351,104.00	1.00	0.00	0.00
BDCP2	Aggregated Change Pointers (BDCP, BDCPS)	13.07	12.82	439,486,947.00	1.00	0.00	0.00
ACCTCR	Compressed Data from FI/CO Document - Currencies	11.92	11.76	347,257,432.00	1.00	0.00	0.00

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Custom Tables

This page provides a focused overview of all custom tables in the SAP S/4HANA system. Custom tables typically follow naming conventions starting with "Z" or "Y", excluding ZARIX tables, which are related to archiving processes and are not considered custom tables. While Partner Code tables (with dedicated namespaces) may qualify as custom tables, they are not included on this page. If such tables are among the 1,000 largest by memory size, they can instead be found on the "Top Tables" page.

Key measures available for analysis on this page include:

- **MEM_GB**: Memory footprint of the table in gigabytes.
- **DISK_GB**: Disk size of the table in gigabytes.
- **PART_COUNT**: Number of partitions in the table (a value of "0" indicates no partitioning).
- **RECORD_COUNT**: Total record count for the table.
- **LOB_DISK_GB**: Size of large objects (LOBs) stored on disk.
- **LOB_MEM_GB**: Size of large objects (LOBs) stored in memory.
- **NSE_GB**: The portion of the table's data moved to NSE (Native Storage Extension), if activated.

This dashboard facilitates in-depth analysis of memory and storage consumption by custom tables, enabling data-driven system optimization and effective capacity planning.

Custom Table Footprint

GB

26

Custom Tables by Memory Size

MEM_GB 1

TABLE_NAME	MEM_GB	DISK_GB	NSE_GB	LOB_DISK_GB	LOB_MEM_GB	PART_COUNT	RECORD_COUNT
ZDMLT_RENUM_FMUS	15.00	22.55	19.67	0.00	0.00	1	909,867,000
ZDMLT_ASLC_DAT	5.82	13.44	11.68	0.00	0.00	1	880,020,195
ZS111	1.19	6.34	0.00	0.00	0.00	1	89,999,996
ZGG02_DOCSEDT	0.37	1.47	0.00	0.00	0.00	1	87,662,916
ZPS_ROOTS_HIER	0.35	2.02	2.00	0.00	0.00	1	72,388,062
ZSWWLOGHIST_TEST	0.34	1.45	0.00	0.00	0.00	5	24,879,712
ZSV_ASSIGNMENT	0.27	2.71	1.82	0.00	0.00	5	128,894,872
ZPMMO_TEST1	0.26	3.68	0.00	0.00	0.00	1	68,686,315
ZZLR_TEST	0.21	5.04	0.00	0.00	0.00	1	52,139,434
ZTD_GEF_MIG_600_root.polygon	0.16	0.17	0.03	0.00	0.00	1	400,000
ZSV1_ASSIGNMENT	0.15	1.44	0.00	0.00	0.00	1	55,913,203
ZKTB_AUAB	0.14	0.99	0.00	0.00	0.00	1	50,000,000
ZFSDM_VDBEPI	0.13	0.73	0.00	0.00	0.00	1	29,999,990
ZFSDM_VDBEPP	0.13	1.24	0.00	0.00	0.00	1	29,999,990
ZKUPMMO_ASSGMT	0.13	1.73	0.00	0.00	0.00	1	55,913,203
ZKKO_ARCHIVED	0.11	0.75	0.75	0.00	0.00	1	17,855,902

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Archiving Achievements

Memory availability can be more accurately determined by leveraging database system statistics during the calculation of archiving metrics. To enable this functionality, please implement SAP Note 2003678, 3144657 & 3314060.

Archiving Achievements

GB

+0

Archiving History

COMPL_ARCH_RUNS 1

ARCH_OBJECT	LAST_RUN	COMPL_ARCH_RUNS	DEL_SIZE_GB	FILE_SIZE_GB
MM_MATNR	13.06.2025	114	0.00	0.00
PS_PROJECT	20.01.2026	110	0.00	0.02
PS_PLAN	12.06.2024	74	0.00	0.00
ZRISKHSTRY	06.10.2020	43	0.00	0.00
CA_BUPA	07.11.2025	40	0.00	0.00
SD_COND	17.07.2023	18	0.00	0.00
PM_ORDER	20.05.2024	12	0.00	0.00
PM_IMRG	14.08.2025	11	0.00	0.00
ZGG02_DOC	08.04.2026	11	0.00	0.27
MM_MATBEL	28.09.2022	10	0.00	0.00
PM_OBJLIST	08.12.2017	10	0.00	0.00
PM_MPLAN	08.04.2025	7	0.00	0.00
RV_LIKP	25.08.2022	7	0.00	0.00
SD_VBAK	25.08.2022	7	0.00	0.00
ZEAM_JBPCK	05.01.2026	7	0.00	0.00
CM_QMEL	20.12.2023	6	0.00	0.00

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Deletion Achievements

This page provides an overview of the deletion jobs executed in the analyzed SAP system. It offers key insights into the execution and status of these jobs, including:

- **PROGNAME:** The name of the deletion program.
- **JOBNAME:** This will provide the job name. The user can schedule a deletion program with different job names.
- **VARIANT:** This will provide the selection criteria used for the deletion job
- **STATUS:** The current status of the program (e.g., completed, in progress, or failed).
- **LAST_EXECUTION:** The date and time of the program's most recent execution.
- **EXECUTION_COUNT:** The total number of times this job has been executed in the system.

Use this page to monitor and analyze deletion activities, ensuring efficient housekeeping and system performance.

Deletion History

2 sorts applied

DOCUMENT_TYPE	PROGNAME	JOBNAME	VARIANT	STATUS	LAST_EXECUTION	EXECUTION_COUNT
Delta Queues	ODQ_CLEANUP	SAP_ODQ_CLEANUP	SAP&24H	Finished/Scheduled	03.07.2026	210
Application Logs	SBAL_DELETE	SAP_REORG_APPLLOG	SAP&S4HREORG	Finished/Scheduled	26.06.2026	144
OTHERS	/LIME/COLLECTION_DELETE	/SCWM/LIME_COLL_DEL	SAP&TECH_JOB	Finished/Scheduled	28.06.2026	75
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_666	&0000000000000	Finished/Scheduled	03.07.2026	15
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_700	&0000000000000	Finished/Scheduled	03.07.2026	15
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_820	&0000000000000	Finished/Scheduled	03.07.2026	15
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_970	&0000000000001	Finished/Scheduled	03.07.2026	15
Shortdumps (Table SNAP)	RSSNAPDL	SAP_REORG_ABAPDUMPS	SAP&001	Finished/Scheduled	03.07.2026	15
Spool Requests	RSPO0041	SAP_REORG_SPOOL	SAP&001	Finished/Scheduled	03.07.2026	15
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_005	&0000000000000	Cancelled	03.07.2026	14
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_515	&0000000000000	Cancelled	03.07.2026	14
Delta Queues	ODQ_CLEANUP	ODQ_CLEANUP_CLIENT_710	&0000000000000	Cancelled	03.07.2026	14
OTHERS	RSBDCREO	SAP_REORG_BATCHINPUT	SAP&001	Finished/Scheduled	03.07.2026	14
Job Logs	RSBTCDEL2	SAP_REORG_MAIL	PSSDEV_MAIL	Finished/Scheduled	02.07.2026	14
Job Logs	RSBTCDEL2	SAP_REORG_JOBS	SAP&001	Finished/Scheduled	03.07.2026	12
OTHERS	RSXMILOGREORG	SAP_REORG_XMILOG	SAP&001	Finished/Scheduled	03.07.2026	3

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Top Growing Tables

This page highlights the Top 100 growing tables in the system, sorted by AVERAGE_MEM_GROWTH_GB, which reflects the average memory growth over the past 30 days. Additional measures provide deeper insights into table growth trends:

- DISK_GROWTH_GB: The increase in disk size over the past 30 days.
- MEM_GROWTH_GB: The increase in memory footprint over the past 30 days.
- RECORD_COUNT_GROWTH: The growth in the number of records during this period.

This data helps identify rapidly growing tables that may not currently be included in any data reduction initiatives. Use this information to proactively address potential storage and memory challenges.

TABLE_NAME

- All
- /AIF/FRUN_LOG
- /AIF/PERS_XML
- /AIF/TRFCQDATA
- /BSNAGT/FILE_INF
- /MFND/D_CLNTSTAT

Apply Selections

Top100 Tables Growth

AVERAGE_MEM_GROWTH_GB

TABLE_NAME	AVERAGE_MEM_GROWTH_GB	MEM_GROWTH_GB	DISK_GROWTH_GB	RECORD_COUNT_GROWTH
ACDOCA	45.6	45.6	+0.0	3,966
MDGD_CP_REP_STAT	41.4	41.4	+0.0	3,209
MOFF	28.3	28.3	0.0	597
FMUSFGA	22.6	22.6	0.0	0
/SCA/KT1SPM	20.5	20.5	0.0	0
/SMB/SBA1_T_FILE	17.6	17.6	0.0	0
AUSP	16.5	16.5	0.0	42
SWWLOGHIST	15.3	15.3	0.0	13,391
SWFREVTCNT	15.0	15.0	0.0	231
ZDMLT_RENUM_FMUS	14.9	14.9	0.0	0
FQM_FLOW	14.4	14.4	-0.4	710
/SAPAPO/DRPMCHNG	13.1	13.1	0.0	0
ACCTCR	11.7	11.7	0.0	5,125
DBTABLOG	11.0	11.0	0.1	426,926
TRGT_USER_DATA	10.9	10.9	0.0	0
EDID4	10.0	10.0	0.0	6,861

SAC Dashboard for Transparency & Monitoring

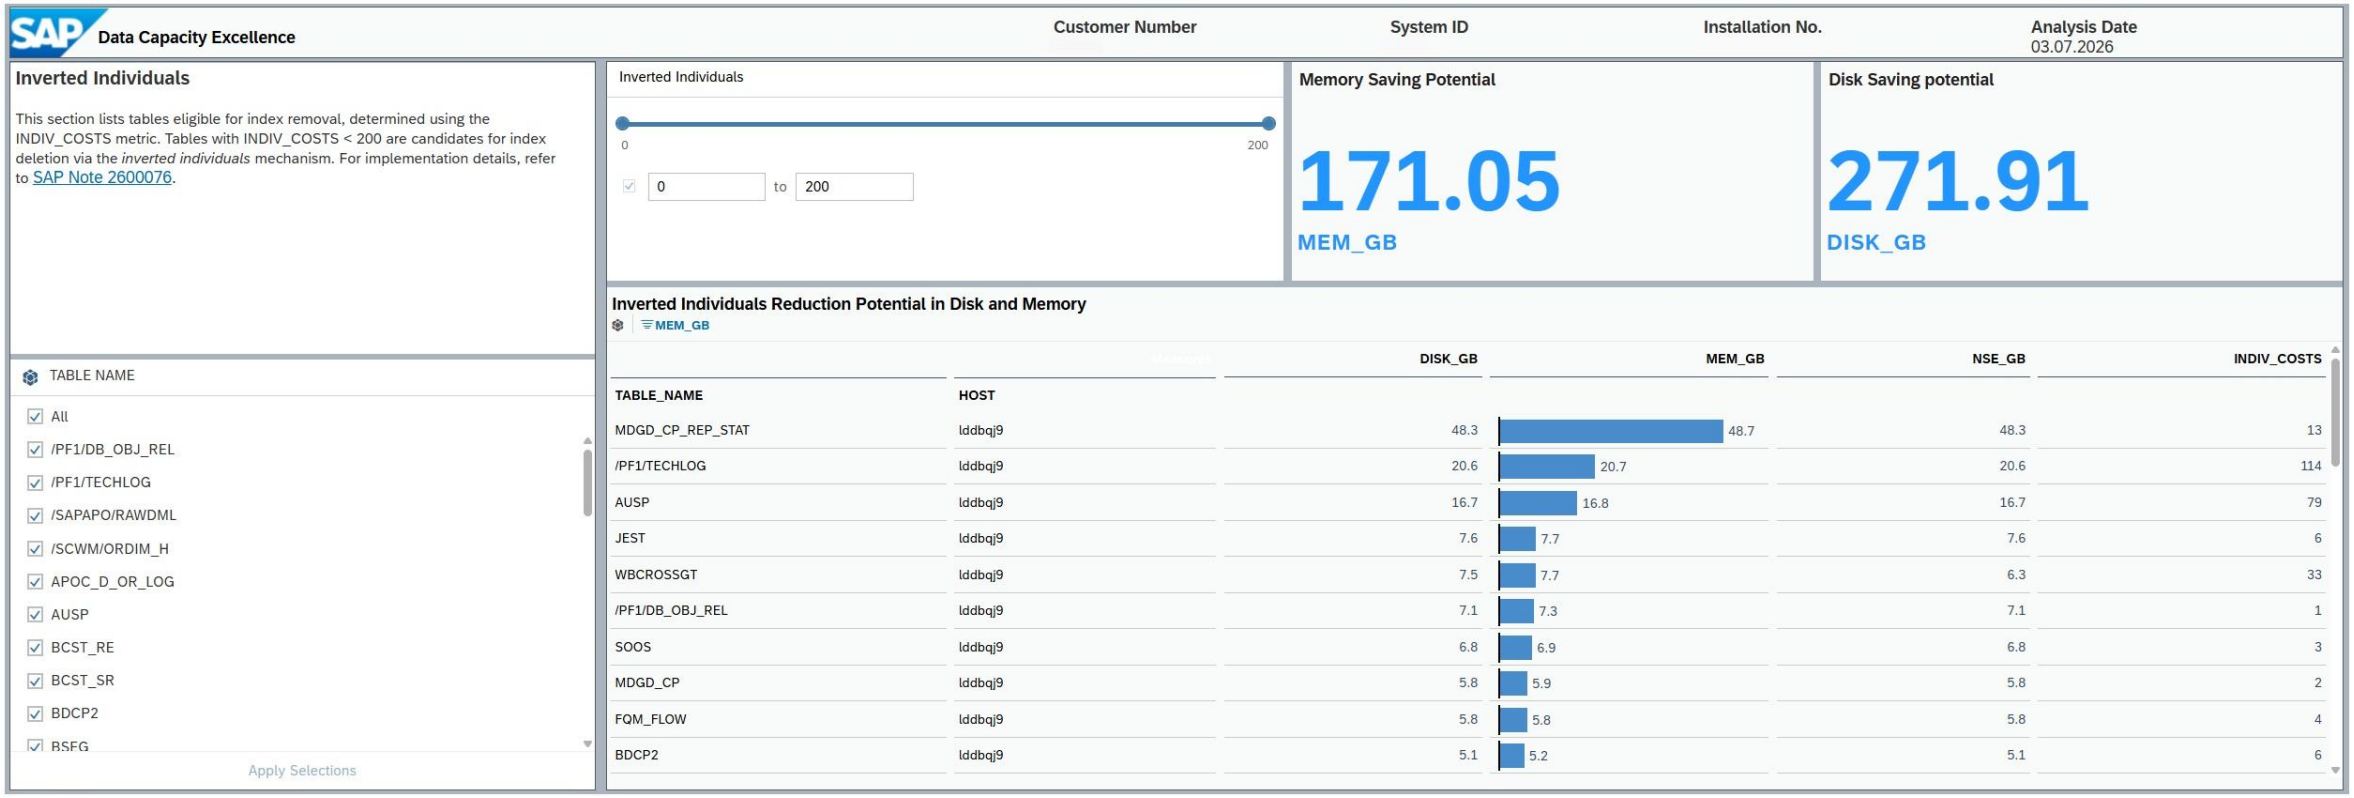

TABLE_NAME	HOST	DISK_GB	MEM_GB	NSE_GB	INDIV_COSTS
MDGD_CP_REP_STAT	lddbqj9	48.3	48.7	48.3	13
/PF1/TECHLOG	lddbqj9	20.6	20.7	20.6	114
AUSP	lddbqj9	16.7	16.8	16.7	79
JEST	lddbqj9	7.6	7.7	7.6	6
WBCROSSGT	lddbqj9	7.5	7.7	6.3	33
/PF1/DB_OBJ_REL	lddbqj9	7.1	7.3	7.1	1
SOOS	lddbqj9	6.8	6.9	6.8	3
MDGD_CP	lddbqj9	5.8	5.9	5.8	2
FQM_FLOW	lddbqj9	5.8	5.8	5.8	4
BDCP2	lddbqj9	5.1	5.2	5.1	6

SAC Dashboard for Transparency & Monitoring

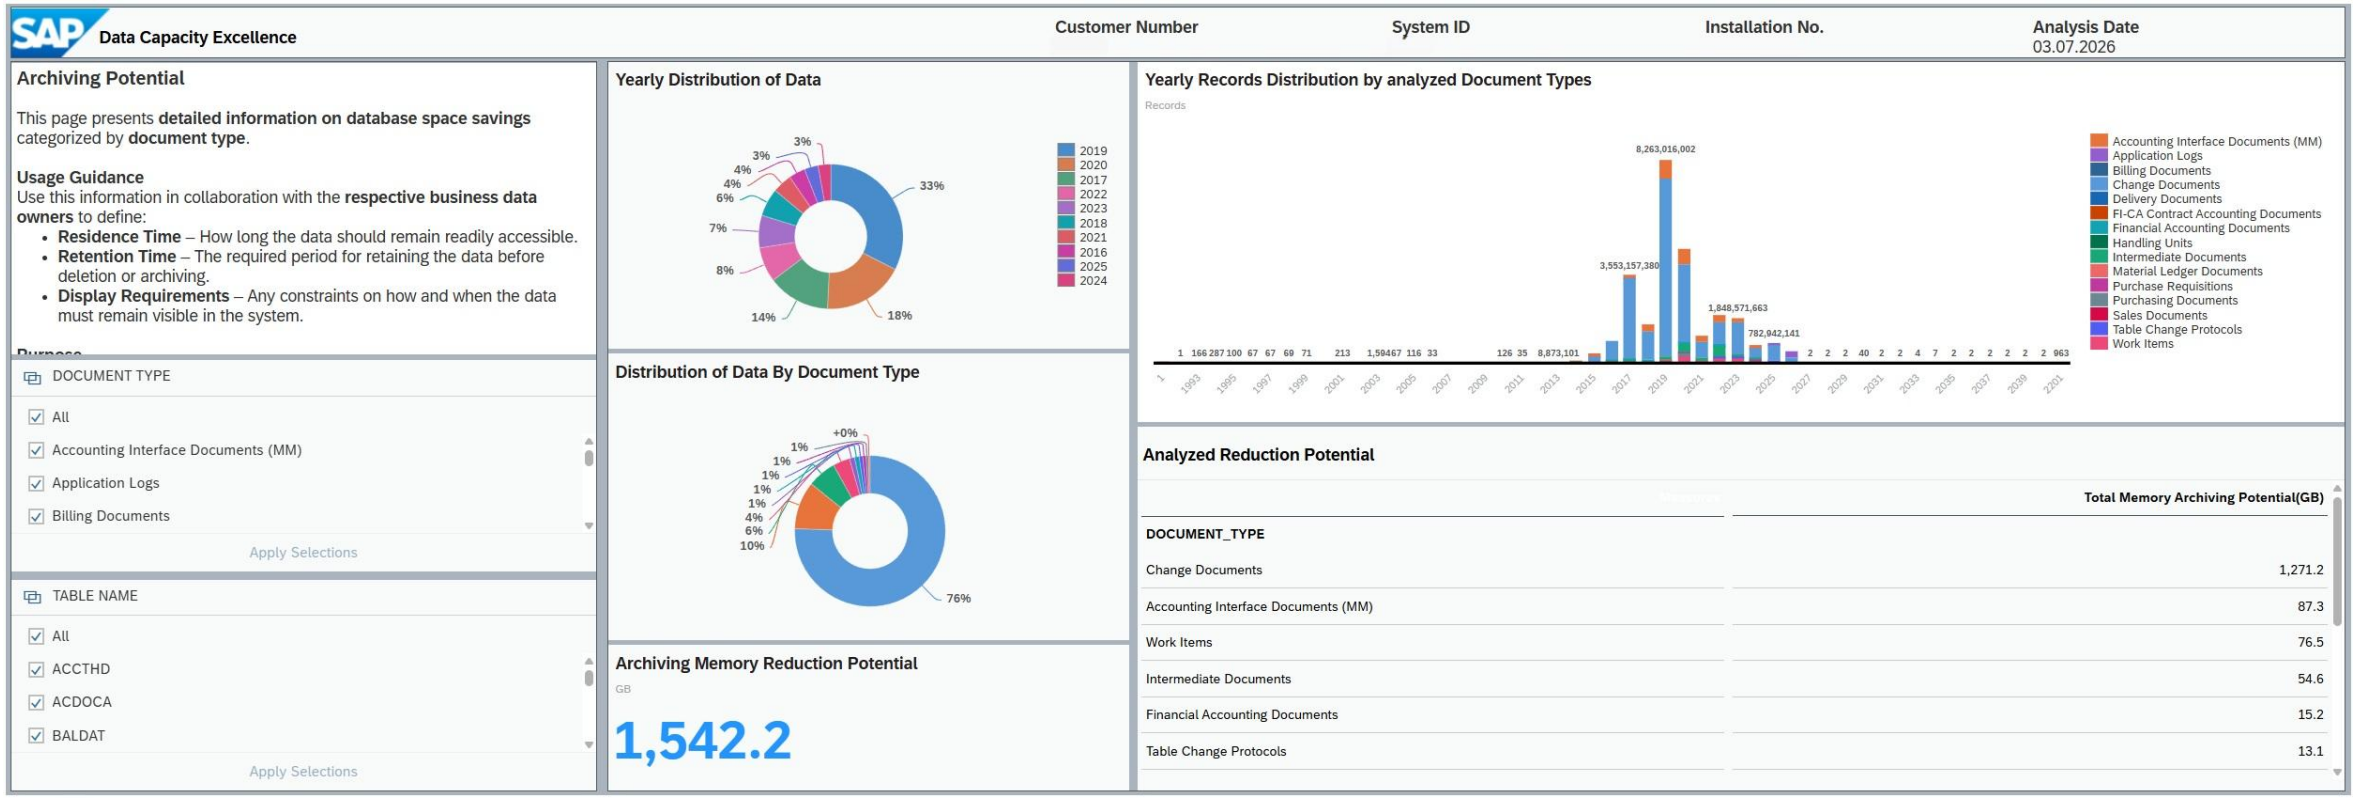

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

NSE

This page provides the tables for which NSE can be activated/deactivated. Tables with Store = 'C', Zero READ_COUNT and Zero WRITE_COUNT are good candidates for NSE activation. Please refer to SAP Note 2799997 for details.

Memory Saving Potential (GB)	System Page Loadable Data Size (GB)	Current Buffer Cache Size (GB)	Out of Buffer Events In the Last 42 days
0.00	4.01	48.40	0

STORE

- All
- C
- R

Apply Selections

TABLE_NAME

- All
- /LIME/NLOG_QUAN
- /SCDL/DB_DF
- /SCWM/ORDIM_C
- /UIF/LREPDCONTCD
- ACCTCR
- ACCTHD
- ACCTIT
- ACDOCA

NSE Reduction Potential in Disk and Memory

2 Filters | READ&WRITE

TABLE_NAME	HOST	STORE	READ_COUNT	WRITE_COUNT	RECORD_COUNT	DISK_GB	MEM_GB	NSE_GB
APQD	vhccuhq4db01	R	0	0	8,862,195	3	2.02	0.00
ECTD_DATA	vhccuhq4db01	C	0	0	310,116	1	0.00	0.00
ECTD_XML_STR	vhccuhq4db01	C	0	0	211,047	9	0.00	0.02
FPLAYOUTT	vhccuhq4db01	R	0	0	14,357	4	0.00	0.00
LTEX	vhccuhq4db01	R	0	0	926,242	3	2.55	0.00
SWNCMONI	vhccuhq4db01	R	0	0	1,846,831	6	5.09	0.00
SXMSCLUR	vhccuhq4db01	R	0	0	1,856,371	7	0.88	0.00
TBTCO	vhccuhq4db01	R	0	0	22,367,814	7	5.42	0.00
TBTCP	vhccuhq4db01	R	0	0	22,367,851	9	7.29	0.00
TST03	vhccuhq4db01	R	0	0	771,846	9	0.05	0.00
ZFIGL079_RAW_CUB	vhccuhq4db01	C	3,095	0	45,067,541	1	0.00	0.00
ZODOO_MOV_INV	vhccuhq4db01	C	3,095	0	1,710,125	1	0.00	0.00
SXMSPHIST	vhccuhq4db01	C	10,537	0	48,432,657	3	0.00	0.00
ZMMIM_MOVI_INVEN	vhccuhq4db01	C	12,436	0	88,414,107	3	0.00	0.00
ZTA_MOVINV058	vhccuhq4db01	C	13,413	0	54,201,284	2	0.00	0.00
ZFIAP_CDGITEM	vhccuhq4db01	C	14,289	0	153,699,549	6	0.00	0.00
ZFIAP_DETALLE	vhccuhq4db01	C	23,290	0	154,162,207	4	0.00	0.00
ZPP_TRAZAB_APL	vhccuhq4db01	C	28,287	2	101,068,222	4	0.00	0.00
ZFI_LOG_INTERFAZ	vhccuhq4db01	C	32,692	0	13,243,769	27	2.00	0.00
EDID4	vhccuhq4db01	C	34,105	5	27,004,701	3	3.16	0.00
ZCO001TT_LEDGER	vhccuhq4db01	C	38,567	0	497,126,893	40	2.43	0.00
MLAUFKEPH	vhccuhq4db01	C	53,507	24	70,992,309	3	2.88	0.00
QASV	vhccuhq4db01	C	54,052	109	58,187,611	1	0.00	0.00
MLPRKEKO	vhccuhq4db01	C	54,160	126	173,310,676	5	3.77	0.00
MLPRKEPH	vhccuhq4db01	C	54,165	126	153,839,469	5	4.09	0.00
SXMSPERROR	vhccuhq4db01	C	54,343	0	295,274	3	0.01	0.00
MLAUFCR	vhccuhq4db01	C	54,444	90	69,482,764	2	2.12	0.00
MLCRF	vhccuhq4db01	C	54,335	202	71,680,351	3	2.79	0.00
/LIME/NLOG_QUAN	vhccuhq4db01	C	53,648	921	16,876,316	2	1.65	0.00

SAC Dashboard for Transparency & Monitoring

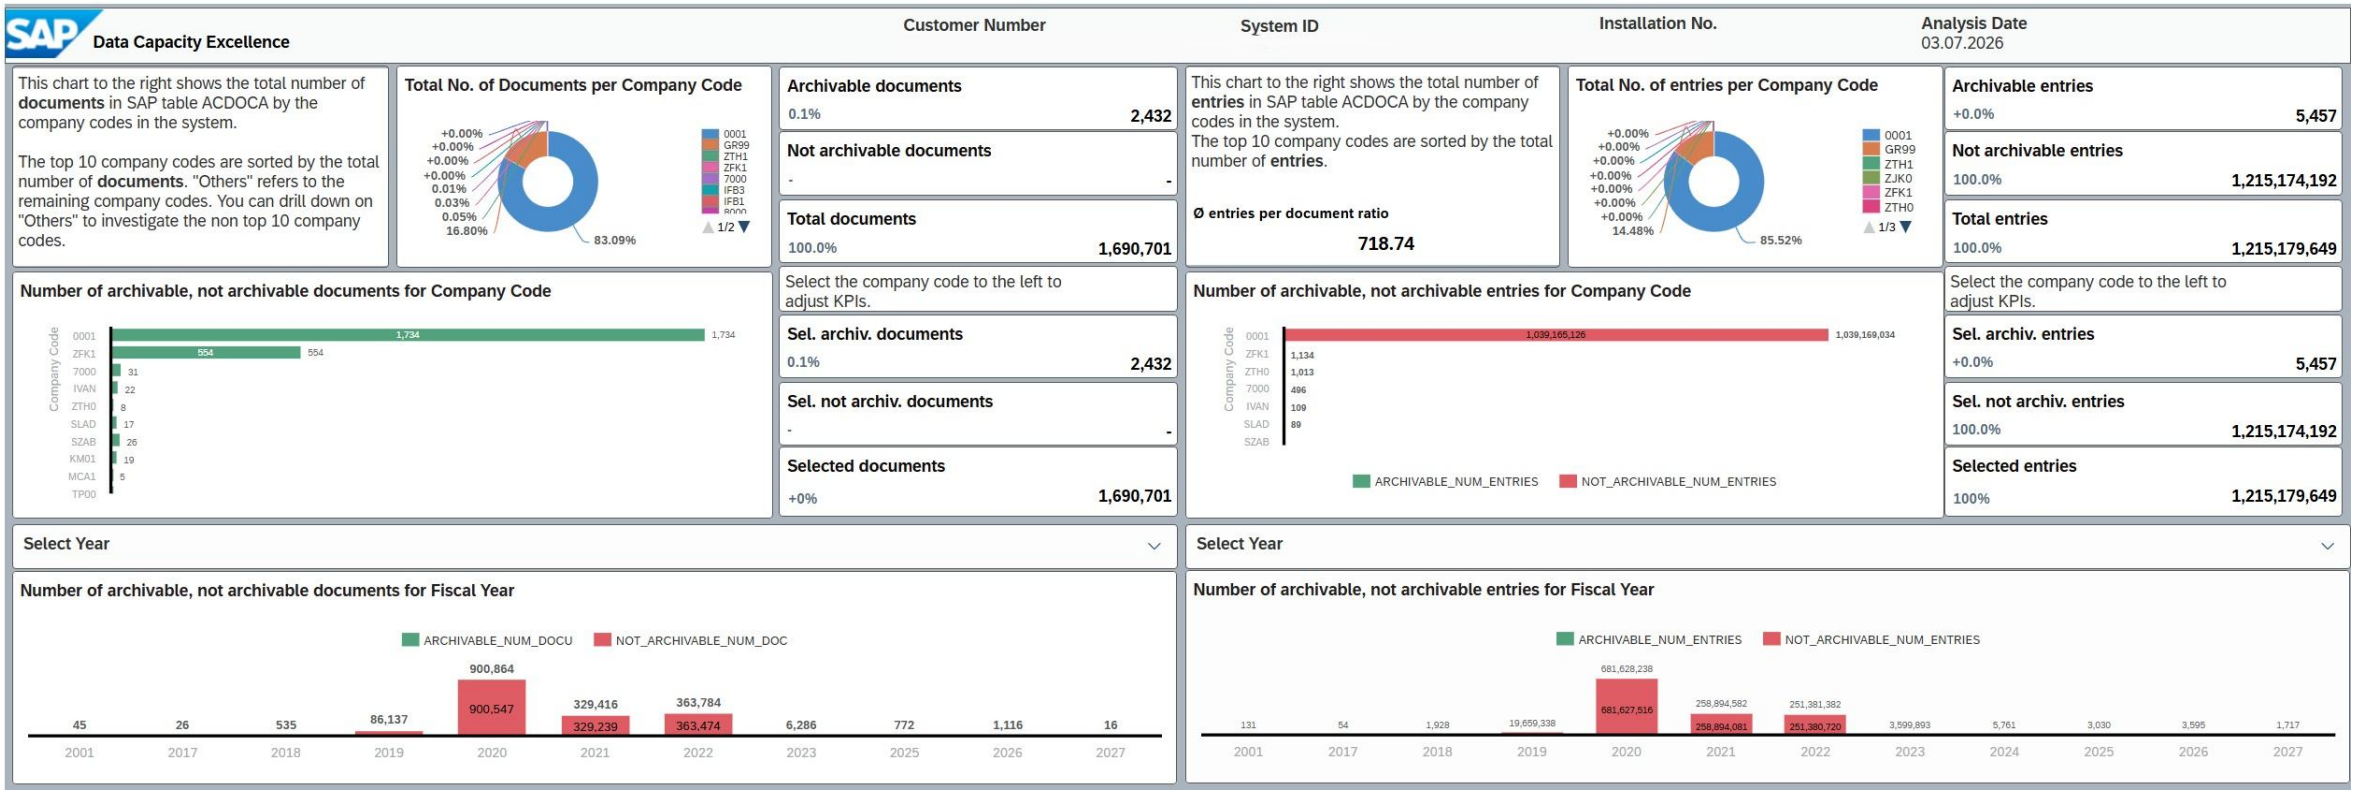

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Tables Logging

This page provides a list of specific tables which are not supposed to be logged in the DBTABLOG table as per the best practice. Table logging in SAP is a mechanism that captures and records all changes made to database table entries. The DBTABLOG table serves as the central repository for storing these change logs, providing a complete audit trail of data modifications including who made the change, when it occurred, and what values were modified.
The list offers insights on the table logging:
•TABNAME: Table Name
•PROTOKOLL: If the value is X then the table is being logged.
•RECORD_COUNT_DBTABLOG: Number of records (Changes) present in the DBTABLOG
•RECORD_COUNT_TOTAL: Total number of records per table

For more information refer to these **SAP notes 2335014 & 2362854.**

PROTOKOLL

- All
- (Null)
- X

Apply Selections

TABLES_LOGGING_DBTABLOG

TABNAME	PROTOKOLL	RECORD_COUNT_DBTABLOG	RECORD_COUNT_TOTAL
/VIRSA/ZFFCDHDR	X	0	0
/VIRSA/ZFFTNSLOG	X	0	0
/VIRSA/ZVIRFFLOG	X	0	0
ALPFASSIGN	(Null)	0	0
AQLTQ	(Null)	5	2,127
AQLTS	(Null)	7,273	39,766
BKK40	(Null)	0	204
BKK42	(Null)	0	221
BKK45	(Null)	0	676
BKK46	(Null)	0	232
BKK98	(Null)	0	0
BKKSONT	X	0	0
COMM_PRPRDCATR	(Null)	0	13
DB6IREORG	(Null)	0	91,054
DB6TREORG	(Null)	0	42,579
DFKKCOLL	(Null)	0	3
DFKKCOLLH	(Null)	0	4
DRFD_REPLICATION	(Null)	0	192,397,032
EKPA	(Null)	0	67,491,717
EVER	(Null)	0	0
FMDECKHK	(Null)	0	0
FMDECKRG	(Null)	0	0
FMDECKRGT	(Null)	0	0
FMDECKUNG	(Null)	0	0
FMSNRULE	(Null)	0	0
FMSNRULEC	(Null)	0	0

SAC Dashboard for Transparency & Monitoring

Data Capacity Excellence

Customer Number	System ID	Installation No.	Analysis Date
			03.07.2026

Tables Partitions

This page provides the list of tables along with its partitions.
TABLE_NAME: Table Name
PART_ID: Partition ID
LEVEL_1: Indicates the partitioning type
LEVEL_2: Indicates the partitioning type Load_Unit: This signifies if the partition is page loadable or column loadable.
RECORD COUNT: The total record count for each partition
DISK_GB: The size of the table on disk, measured in gigabytes.
NSE_GB: The size of the table on NSE, measured in gigabytes
MEM_GB: The memory footprint of each table, measured in gigabytes
BUFFER_GB: Indicates the buffer cache used for loading.

For more information refer to these **SAP notes 2044468**

TABLE_NAME

- [ ] ACDOCA

Apply Selections

LOAD_UNIT

- [ ] COLUMN

Apply Selections

Record Count in Table/Partitions

0 — 2147483648

0 to 2147483648

Table Partitions

TABLE_NAME	PART_ID	LEVEL_1	LEVEL_2	LOAD_UNIT	BUFFER_GB	DISK_GB	MEM_GB	NSE_GB	RECORD_COUNT
ACDOCA	1	1 (HASH)	-	COLUMN	0	22	22		674,007,321
ACDOCA	2	2 (HASH)	-	COLUMN	0	22	22		674,019,680
ACDOCA	3	3 (HASH)	-	COLUMN	0	22	22		676,180,259
ACDOCA	4	4 (HASH)	-	COLUMN	0	22	22		674,623,888

Key Take Aways

What is it ?

An automated , AI powered product built into SAP Cloud ALM that continuously monitor your SAP HANA database health.

What it does ?

Finds data that is costing you money and gives you step by step actions to fix

Who it’s for ?

Any SAP customer on S/4 HANA with Advance or Max success plan

Q&A

Thank you.

Pankaj Acharya
DCE-Product Owner
pankaj.acharya@sap.com

© 2025 SAP SE or an SAP affiliate company. All rights reserved. See Legal Notice on www.sap.com/legal-notice for use terms, disclaimers, disclosures, or restrictions related to this material.